

Rust out your C

Carol (Nichols || Goulding)
@carols10cents



Agenda

- 
- 1 **Caveats**
 - 2 **Background**
 - 3 **Techniques**
 - 4 **Tricky Stuff**
 - 5 **Benefits?**



DO NOT

Bad reasons to rewrite your C in Rust

- Cool kids
- I feel like it
- I'm bored
- Job security
- Carol said

Good reasons to rewrite your C in Rust

- Performance
- Safety
- Lower maintenance costs
- Expand # of maintainers
- For fun, not work

I FIGHT FOR
THE USERS

Any time you
rewrite,
code will
get better.

Things I knew before

- Rust
- Legacy code
- Testing

Things I
DID NOT
know before

- C
- FFI
- zopfli

Background: zopfli



CODING HORROR

programming and human factors

Google™ Custom Search



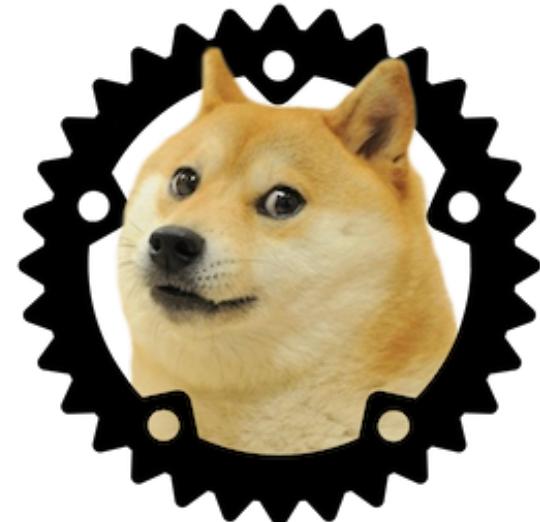
02 Jan 2016

Zopfli Optimization: Literally Free Bandwidth

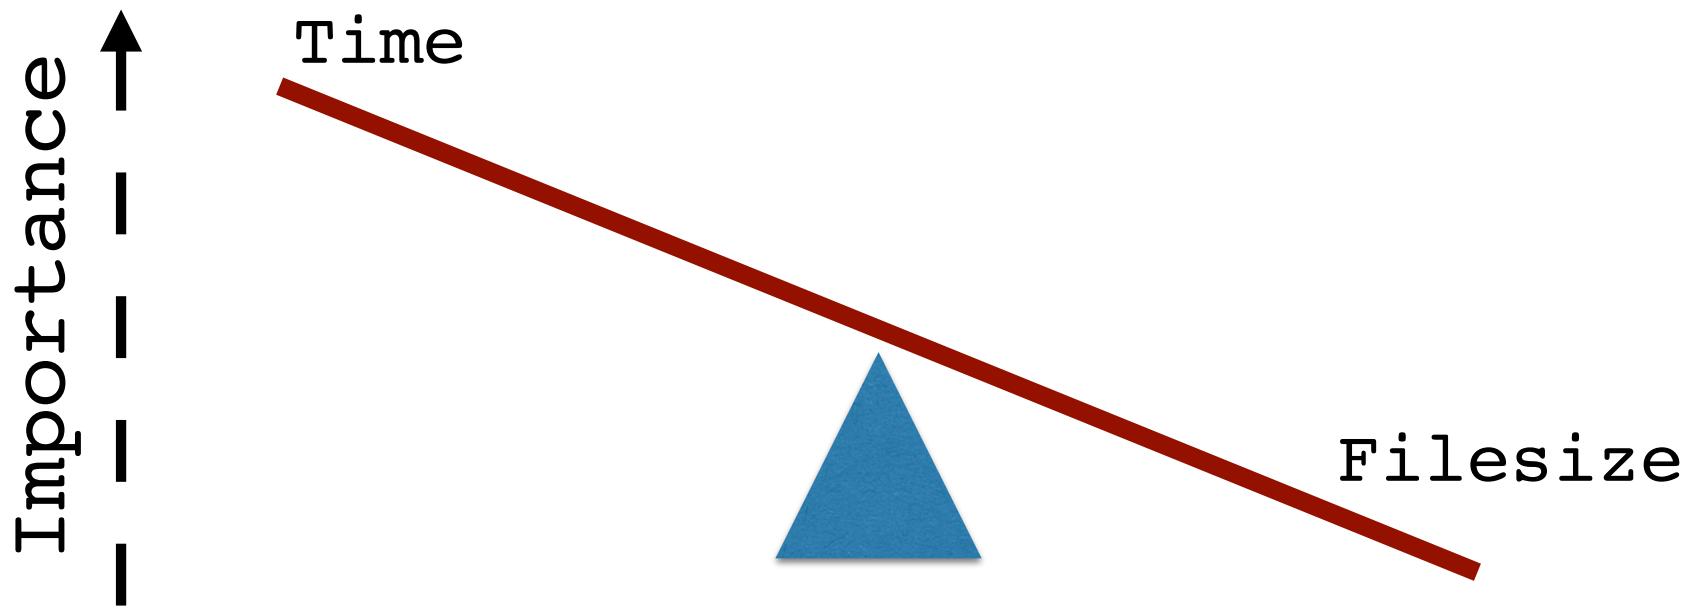
11

<https://blog.codinghorror.com/zopfli-optimization-literally-free-bandwidth/>

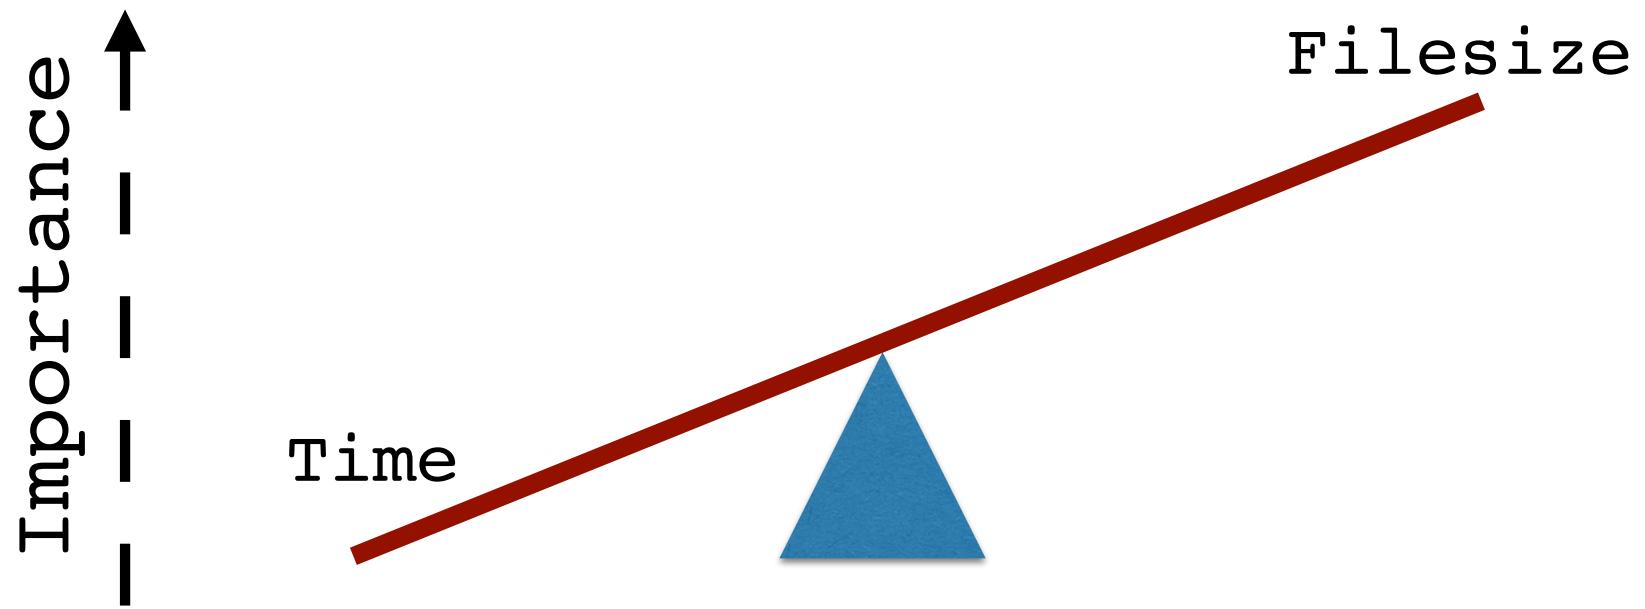
This is a DEFLATE
file, I know this!



gzip



zopfli





eliotsykes commented on Dec 8, 2015



The above benchmark is with Zopfli using its default 15 iterations. Results with a single Zopfli iteration bring the difference down to zlib being ~25 times faster than Zopfli. A single Zopfli iteration is almost as good as 15 Zopfli iterations in terms of the sample file size reductions.

```
Calculating -----
      zlib      9.000  i/100ms
zopfli (1 iteration)    1.000  i/100ms
-----
      zlib     96.592  (± 6.2%) i/s -    486.000
zopfli (1 iteration)    4.195  (± 0.0%) i/s -    21.000
```



schnneems commented on Dec 8, 2015

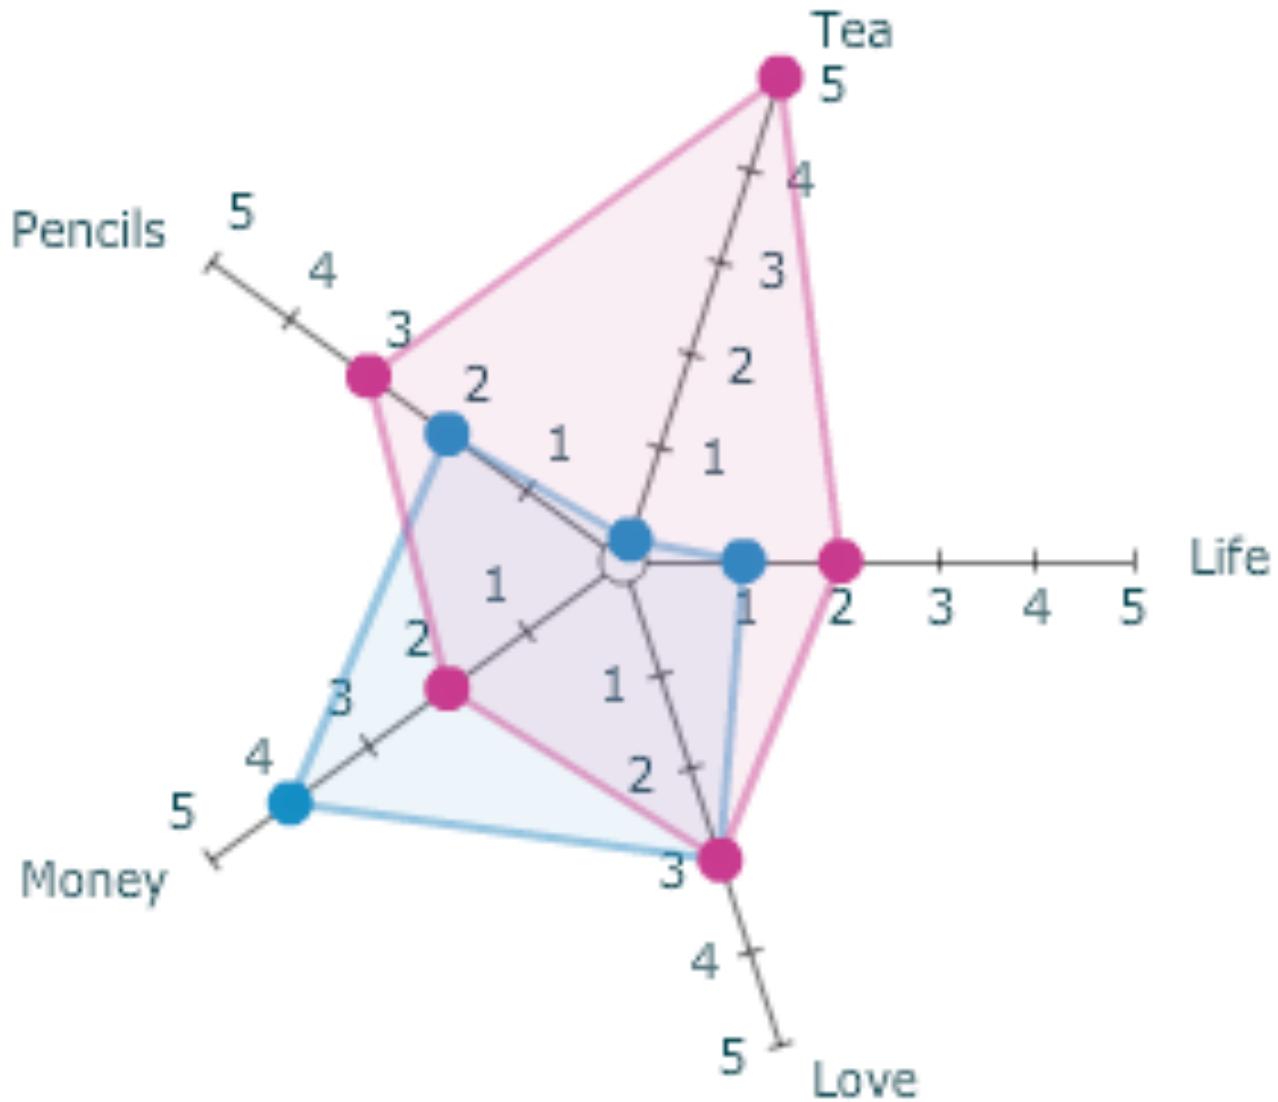
Ruby on Rails member



That's better, still a bit too slow to make the default I think. Maybe we can add it in and make it configurable.

It looks like very little of the code is actually written in C. We could probably get a larger speed up by re-writing more of it in C and doing some benchmarking.

I have another concern with adding this in. I know libraries like Rails ship with gems with C extensions (nokogiri) and somehow they manage to play nice with other rubies like JRuby, but i'm not sure how exactly. We need to make sure we don't break jruby compatibility. There's no way to conditionally add something to a gemspec based on Ruby implementation (that i'm aware of). I also want to be cautious about adding c-extensions to dependencies. They take much longer to install, and by declaring it in the gemspec it would be installed even if someone was not using it. Deploy timeouts from too many c-extensions are a thing.



example radar chart demoing MIT Licensed flex component from
<http://www.boost.co.nz/blog/2009/05/flex-radar-chart-component/>

Techniques

IF IT AIN'T BAROQUE



DON'T FIX IT

Golden Master Tests

Makefile changes

```
+ZOPFLI_RUST_RELEASE := target/release/libzopfli.a  
ZOPFLILIB_SRC = src/zopfli/blocksplitter.c src/zopfli/cache.c...
```

Makefile changes

```
ZOPFLI_RUST_RELEASE := target/release/libzopfli.a
ZOPFLILIB_SRC = src/zopfli/blocksplitter.c src/zopfli/cache.c...
+.PHONY: target/release/libzopfli.a
+target/release/libzopfli.a:
+    cargo build --verbose --release
```

Makefile changes

```
ZOPFLI_RUST_RELEASE := target/release/libzopfli.a
ZOPFLILIB_SRC = src/zopfli/blocksplitter.c src/zopfli/cache.c...
.PHONY: target/release/libzopfli.a
target/release/libzopfli.a:
    cargo build --verbose --release

-zopfli:
-    $(CC) $(ZOPFLILIB_SRC) $(ZOPFLIBIN_SRC) $(CFLAGS) -o zopfli
+zopfli: $(ZOPFLI_RUST_RELEASE)
+    $(CC) $(ZOPFLI_RUST_RELEASE) $(ZOPFLILIB_SRC) $(ZOPFLIBIN_SRC) $(CFLAGS) -o zopfli
```

Cargo.toml changes

```
+ [lib]
+crate-type = ["staticlib"]
```

Remove a function from C

```
size_t CalculateTreeSize(const unsigned* ll_lengths, const unsigned* d_lengths) {  
    size_t result = 0;  
    int i;  
  
    for(i = 0; i < 8; i++) {  
        size_t size = EncodeTreeNoOutput(ll_lengths, d_lengths,  
                                         i & 1, i & 2, i & 4);  
        if (result == 0 || size < result) result = size;  
    }  
  
    return result;  
}
```

Remove a function from C

```
extern size_t CalculateTreeSize(const unsigned* ll_lengths, const unsigned*  
d_lengths);
```

Add a function to Rust

```
size_t CalculateTreeSize(const unsigned* ll_lengths, const unsigned* d_lengths) {
    size_t result = 0;
    int i;

    for(i = 0; i < 8; i++) {
        size_t size = EncodeTreeNoOutput(ll_lengths, d_lengths,
                                         i & 1, i & 2, i & 4);
        if (result == 0 || size < result) result = size;
    }

    return result;
}
```

Add a function to Rust

```
#[no_mangle]
#[allow(non_snake_case)]
size_t CalculateTreeSize(const unsigned* ll_lengths, const unsigned* d_lengths) {
    size_t result = 0;
    int i;

    for(i = 0; i < 8; i++) {
        size_t size = EncodeTreeNo0Output(ll_lengths, d_lengths,
                                         i & 1, i & 2, i & 4);
        if (result == 0 || size < result) result = size;
    }

    return result;
}
```

Add a function to Rust

```
#[no_mangle]
#[allow(non_snake_case)]
pub extern fn CalculateTreeSize(const unsigned* ll_lengths, const unsigned* d_lengths) -> size_t {
    size_t result = 0;
    int i;

    for(i = 0; i < 8; i++) {
        size_t size = EncodeTreeNoOutput(ll_lengths, d_lengths,
                                         i & 1, i & 2, i & 4);
        if (result == 0 || size < result) result = size;
    }

    return result;
}
```

Add a function to Rust

```
#[no_mangle]
#[allow(non_snake_case)]
pub extern fn CalculateTreeSize(ll_lengths: const unsigned*, d_lengths: const
unsigned*) -> size_t {
    size_t result = 0;
    int i;

    for(i = 0; i < 8; i++) {
        size_t size = EncodeTreeNoOutput(ll_lengths, d_lengths,
                                         i & 1, i & 2, i & 4);
        if (result == 0 || size < result) result = size;
    }

    return result;
}
```

Add a function to Rust

```
use libc::{size_t, c_uint};

#[no_mangle]
#[allow(non_snake_case)]
pub extern fn CalculateTreeSize(ll_lengths: *const c_uint, d_lengths: *const
c_uint) -> size_t {
    size_t result = 0;
    int i;

    for(i = 0; i < 8; i++) {
        size_t size = EncodeTreeNo0output(ll_lengths, d_lengths,
                                         i & 1, i & 2, i & 4);
        if (result == 0 || size < result) result = size;
    }

    return result;
}
```

Add a function to Rust

```
use libc::{size_t, c_uint};

#[no_mangle]
#[allow(non_snake_case)]
pub extern fn CalculateTreeSize(ll_lengths: *const c_uint, d_lengths: *const
c_uint) -> size_t {
    let mut result = 0;
    int i;

    for(i = 0; i < 8; i++) {
        let size = EncodeTreeNoOutput(ll_lengths, d_lengths,
                                      i & 1, i & 2, i & 4);
        if (result == 0 || size < result) result = size;
    }

    return result;
}
```

Add a function to Rust

```
use libc::{size_t, c_uint};

#[no_mangle]
#[allow(non_snake_case)]
pub extern fn CalculateTreeSize(ll_lengths: *const c_uint, d_lengths: *const
c_uint) -> size_t {
    let mut result = 0;

    for i in 0..8 {
        let size = EncodeTreeNoOutput(ll_lengths, d_lengths,
                                      i & 1, i & 2, i & 4);
        if (result == 0 || size < result) result = size;
    }

    return result;
}
```

Add a function to Rust

```
use libc::{size_t, c_uint};

#[no_mangle]
#[allow(non_snake_case)]
pub extern fn CalculateTreeSize(ll_lengths: *const c_uint, d_lengths: *const
c_uint) -> size_t {
    let mut result = 0;

    for i in 0..8 {
        let size = EncodeTreeNoOutput(ll_lengths, d_lengths,
                                      i & 1, i & 2, i & 4);
        if result == 0 || size < result {
            result = size;
        }
    }

    return result;
}
```

Add a function to Rust

```
use libc::{size_t, c_uint};

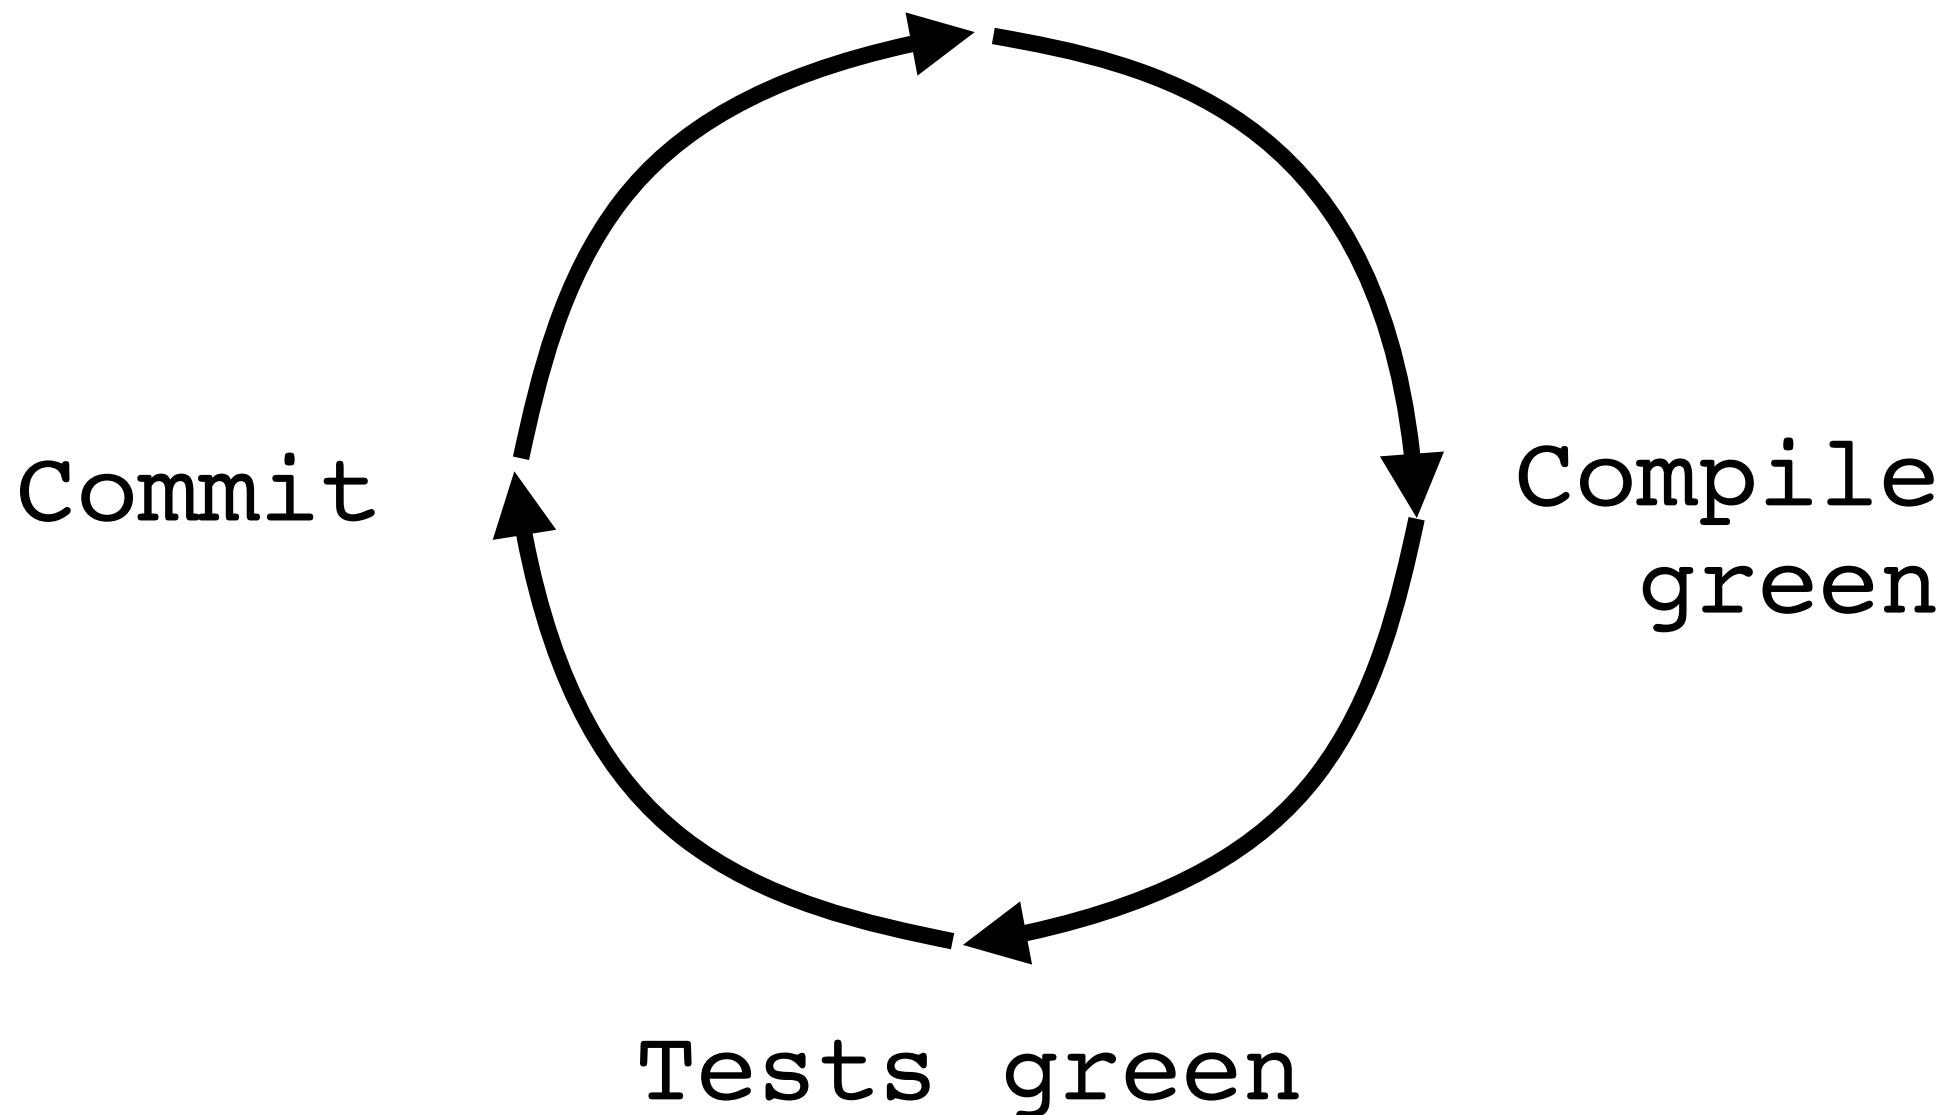
#[no_mangle]
#[allow(non_snake_case)]
pub extern fn CalculateTreeSize(ll_lengths: *const c_uint, d_lengths: *const
c_uint) -> size_t {
    let mut result = 0;

    for i in 0..8 {
        let size = EncodeTreeNoOutput(ll_lengths, d_lengths,
                                      i & 1, i & 2, i & 4);
        if result == 0 || size < result {
            result = size;
        }
    }

    result
}
```

Incremental

Move C to Rust



“Are you
done yet?”

BRING ME

SCHRÖDINGER'S HEAD

What if it
doesn't pass
the tests?

git checkout!
take a smaller
step.

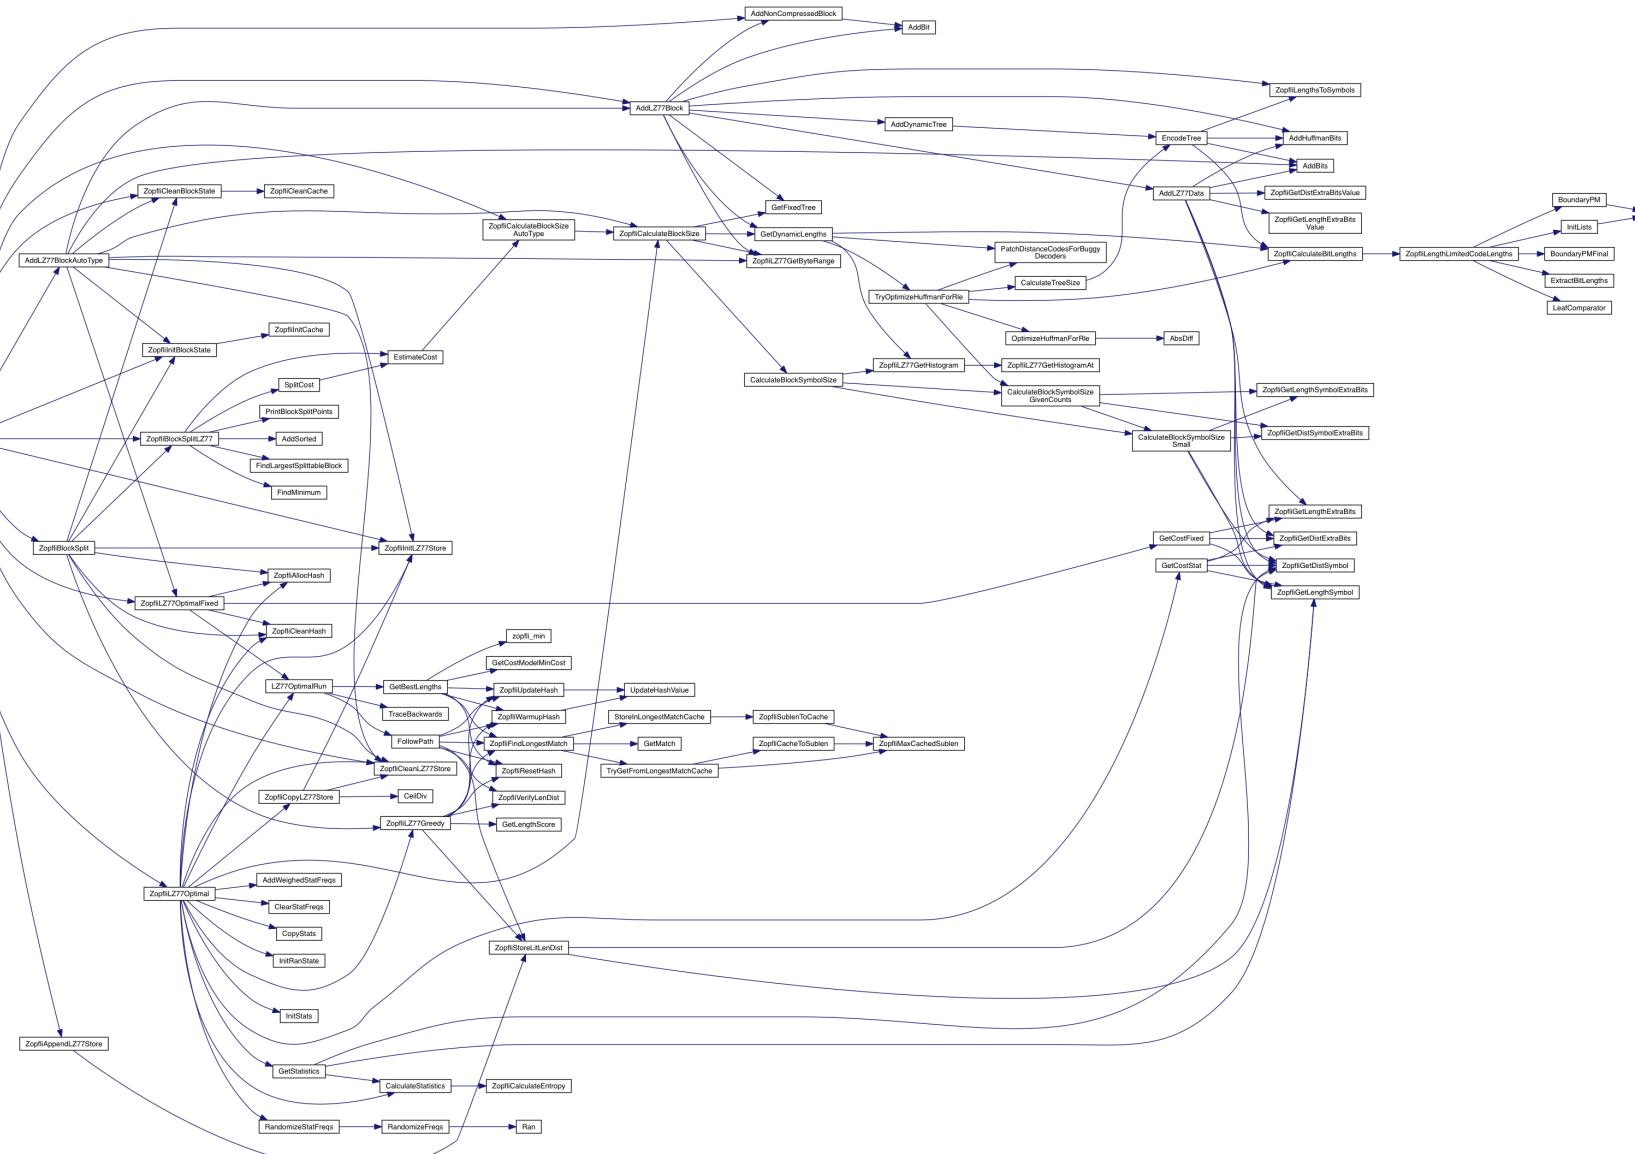
smaller

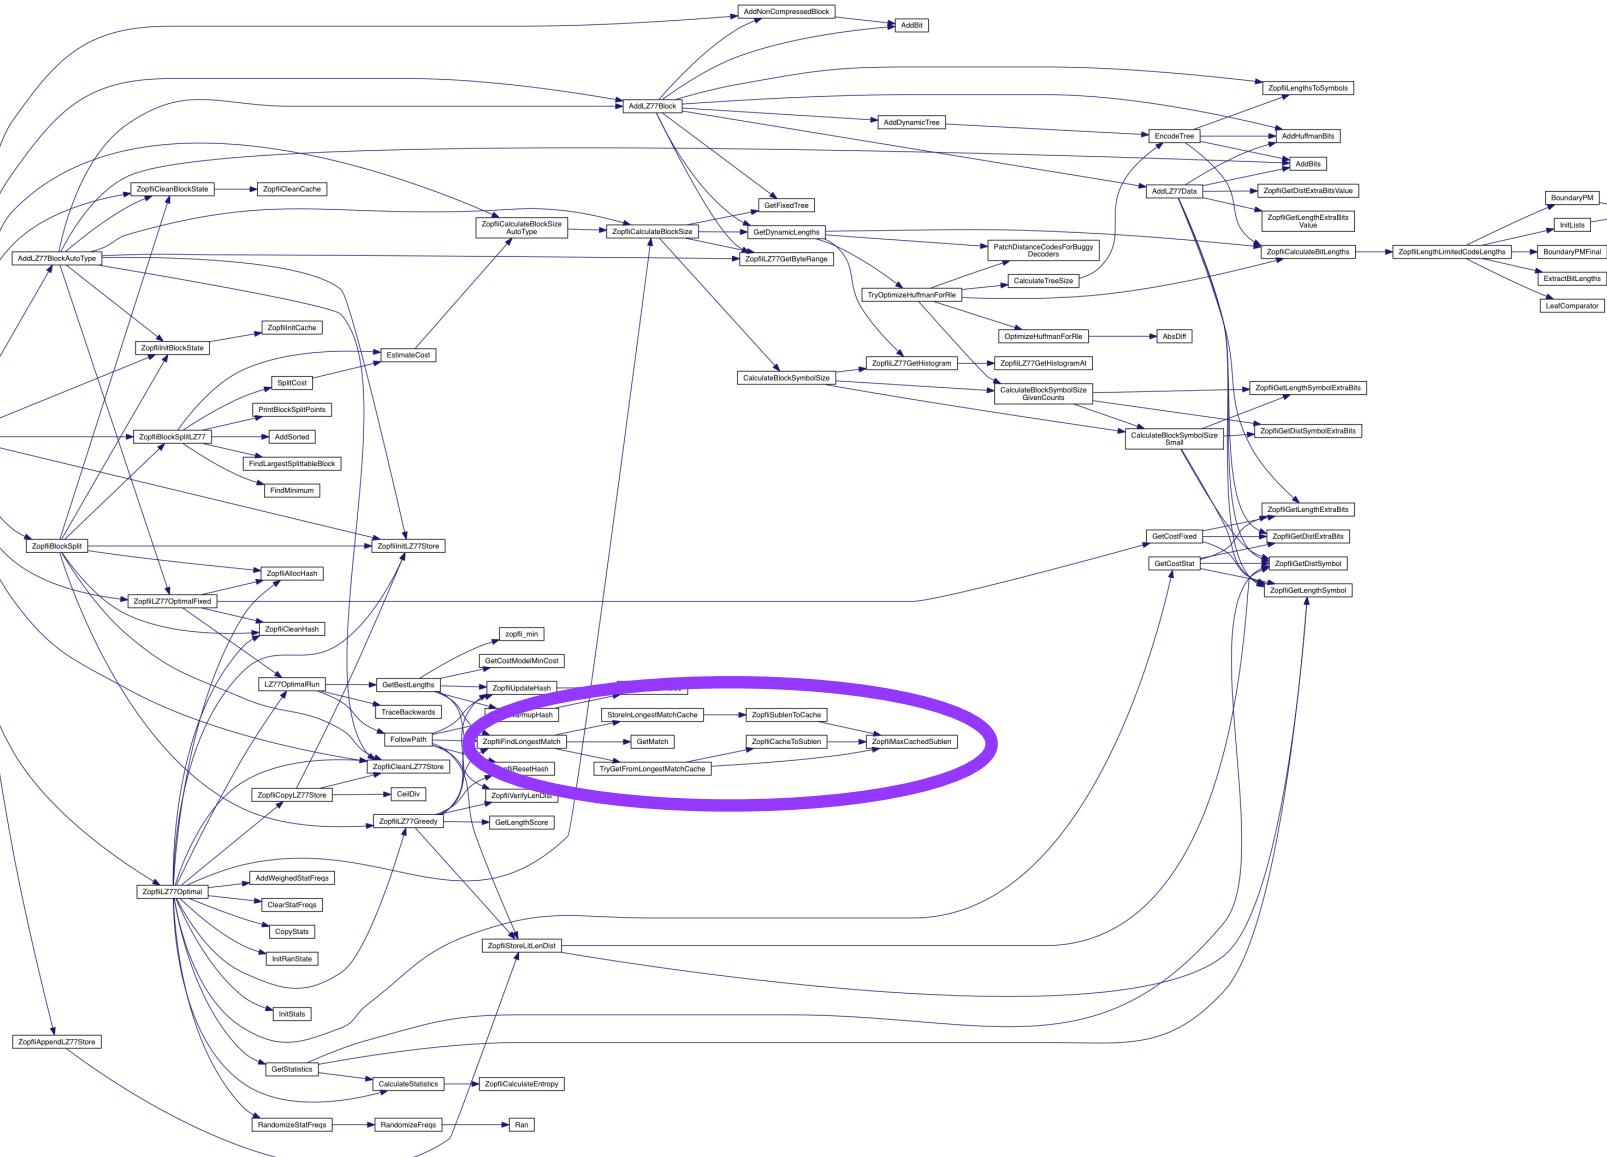
- Extract functions
- Make the C more like Rust first
- Don't make the Rust idiomatic at all, even when it seems easy

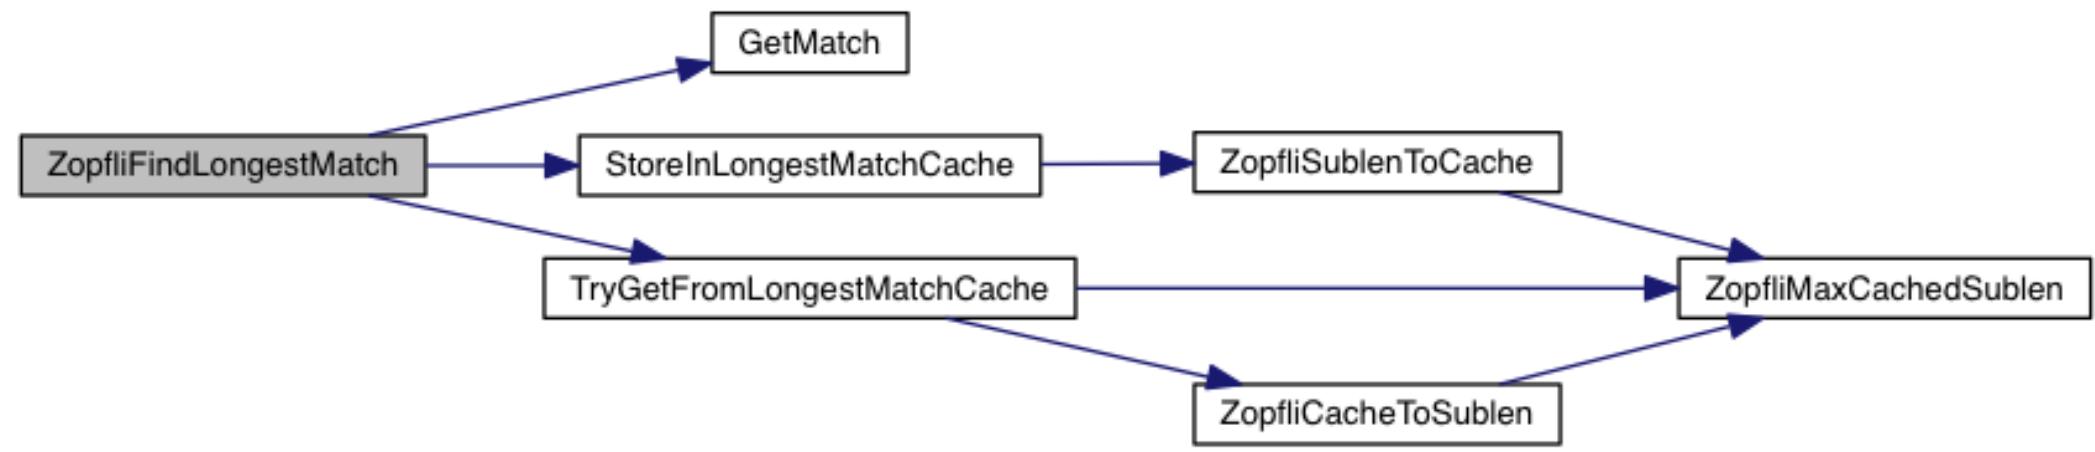
 how

what

Follow the
call graph



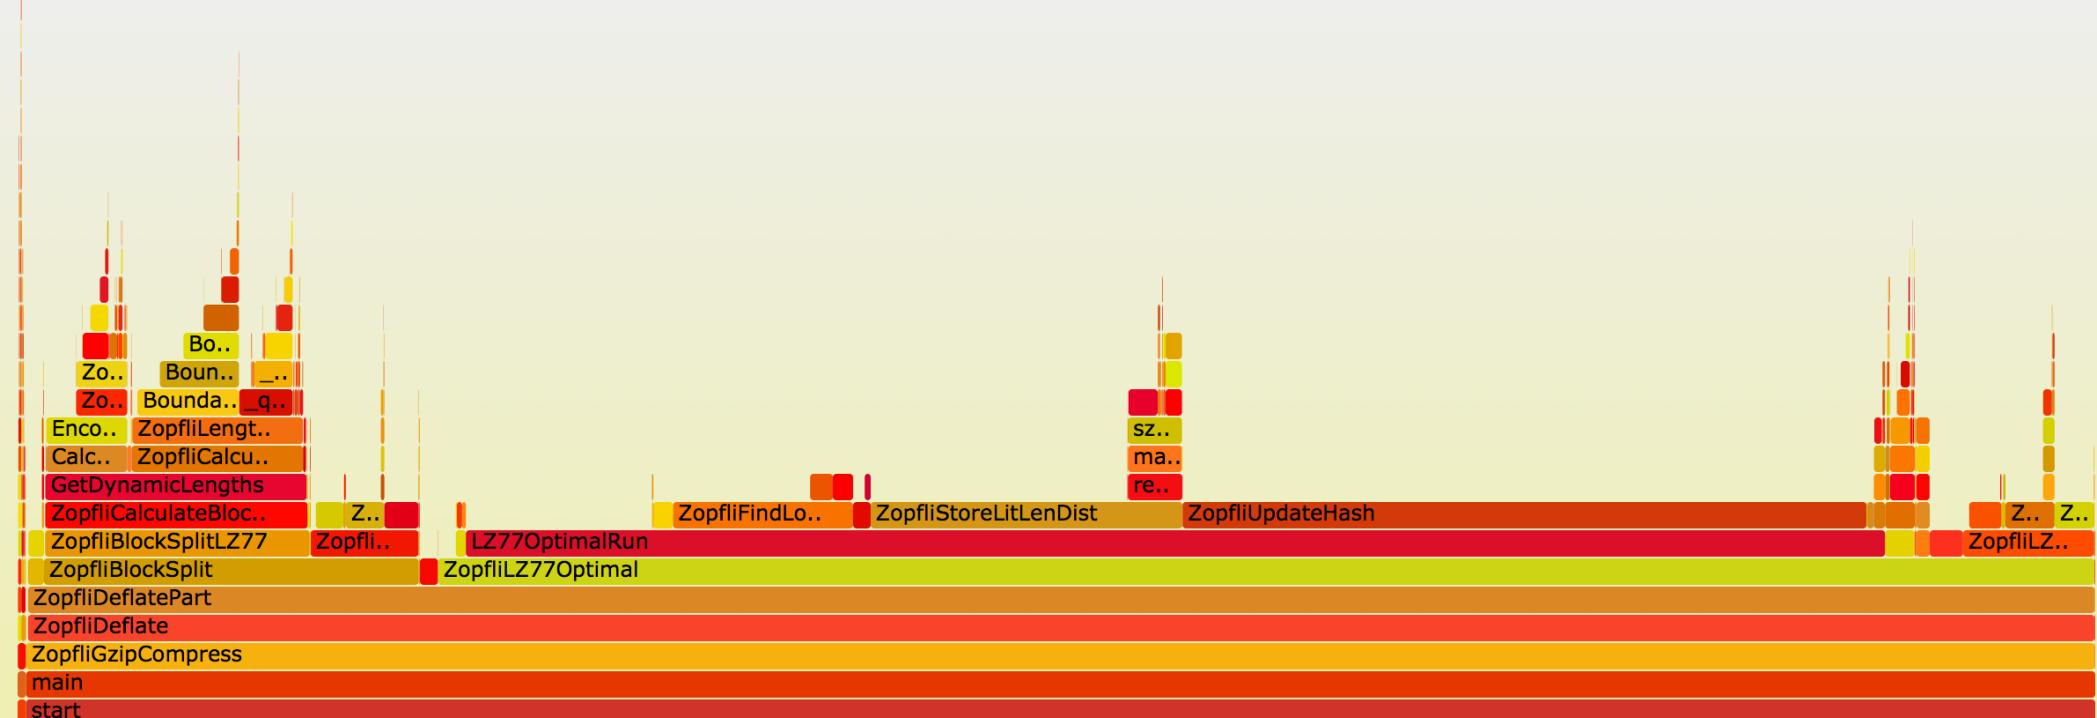




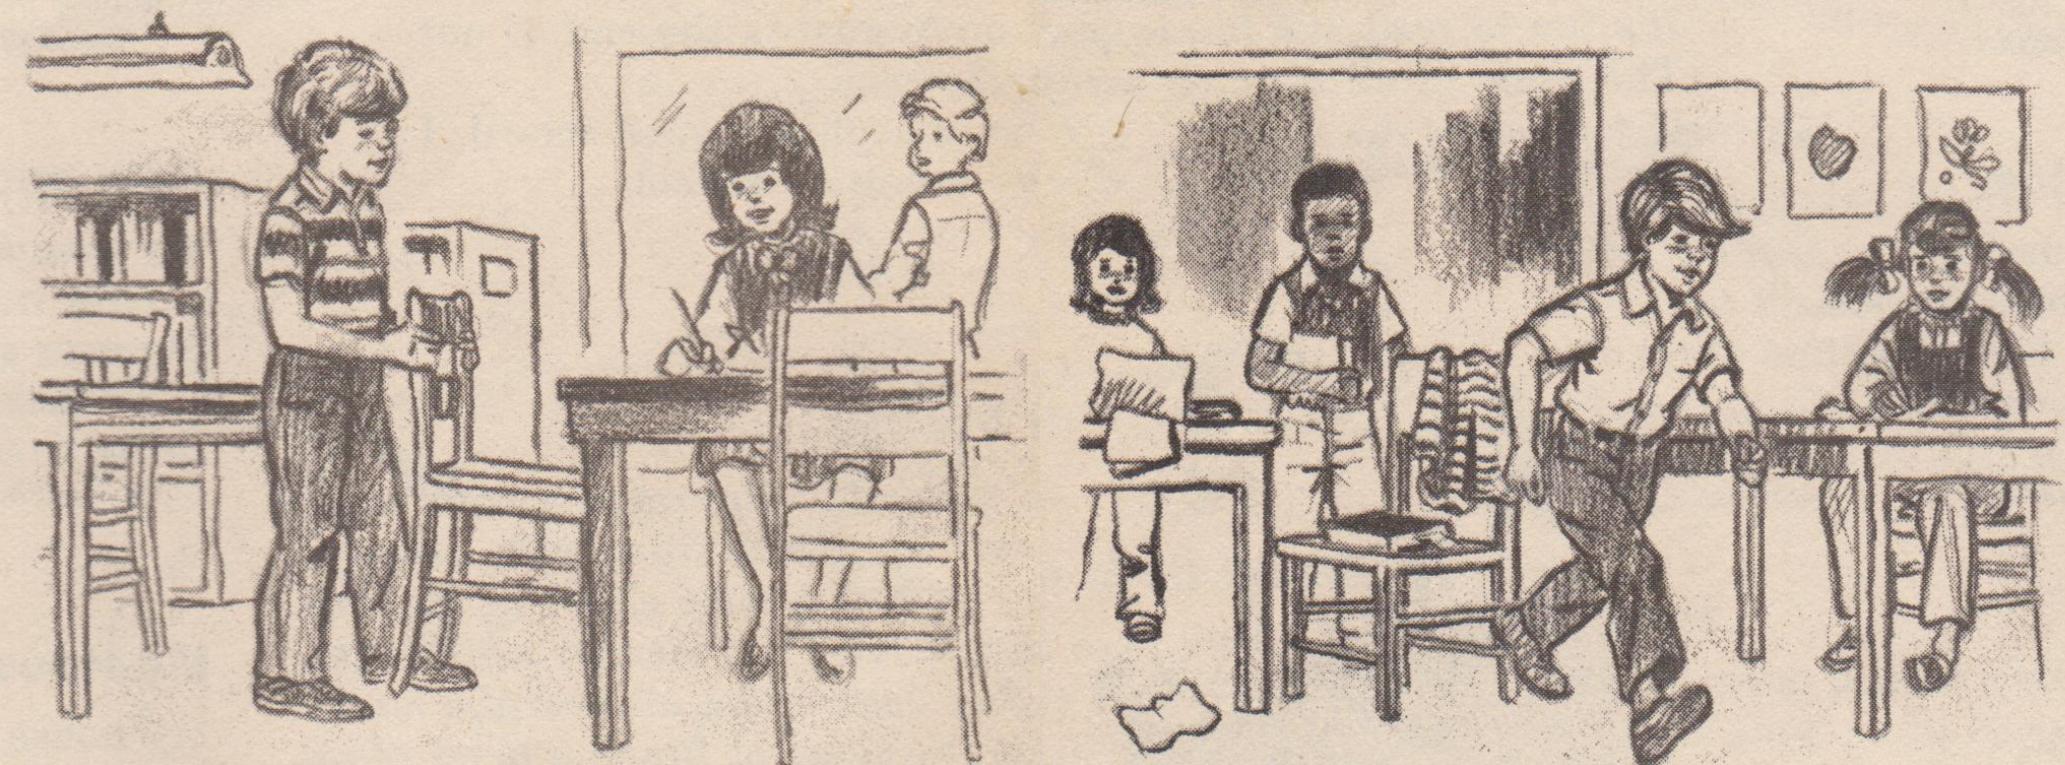
Profile

Flame Graph

Search



Tricky
stuff



Rust upholds your invariants.

C says, “Nuts to your invariants.”

```
static void RandomizeFreqs(RanState*  
state, size_t* freqs, int n) {  
  
    int i;  
    for (i = 0; i < n; i++) {  
  
        if ((Ran(state) >> 4) % 3 == 0) {  
            freqs[i] = freqs[Ran(state) % n];  
        }  
    }  
}
```

```
static void RandomizeFreqs(RanState*  
state, size_t* freqs, int n) {  
  
    int i;  
    for (i = 0; i < n; i++) {  
  
        if (((Ran(state) >> 4) % 3 == 0) {  
            freqs[i] = freqs[Ran(state) % n];  
        }  
    }  
}
```

```
for freq in freqs.iter_mut() {  
    if (Ran(state) >> 4) % 3 == 0 {  
        let index = (Ran(state) % n as  
u32) as usize;  
        *freq = freqs[index];  
    }  
}
```

```
for freq in freqs.iter_mut() {  
    if (Ran(state) >> 4) % 3 == 0 {  
        let index = (Ran(state)) % n as  
u32) as usize;  
        *freq = freqs[index];  
    }  
}
```

error: cannot use `freqs[..]` because it
was mutably borrowed [E0503]

```
*freq = freqs[index];  
          ^~~~~~
```

note: borrow of `*freqs` occurs here
for freq in freqs.iter_mut() {
 ^~~~

```
for i in 0..n {  
    if (Ran(state) >> 4) % 3 == 0 {  
        let index = (Ran(state) % n as  
u32) as usize;  
        freqs[i] = freqs[index];  
    }  
}
```

RandomizeFreqs...
shuffle?

Type conversions



```
double GetBestLengths(..., float* costs, ...) {  
  
    double newCost = costmodel(in[i]) + costs[j];  
  
    if (newCost < costs[j + 1]) {  
        costs[j + 1] = newCost;  
    }  
    ...  
}
```

Type conversions



```
double GetBestLengths(..., float* costs, ...) {  
  
    double newCost = costmodel(in[i]) + costs[j];  
  
    if (newCost < costs[j + 1]) {  
        costs[j + 1] = newCost;  
    }  
    ...  
}
```

```
fn GetBestLengths(..., costs: *mut c_float, ...) -> c_double {  
  
    let newCost = costmodel(in[i]) + costs[j] as c_double;  
  
    if newCost < costs[j + 1] as c_double {  
        costs[j + 1] = newCost as c_float;  
    }  
    ...  
}
```

Type conversions



```
double GetBestLengths(..., float* costs, ...) {  
  
    double newCost = costmodel(in[i]) + costs[j];  
  
    if (newCost < costs[j + 1]) {  
        costs[j + 1] = newCost;  
    }  
    ...  
}
```

```
fn GetBestLengths(..., costs: *mut c_float, ...) -> c_double {  
  
    let newCost = costmodel(in[i]) + costs[j] as c_double;  
  
    if newCost < costs[j + 1] as c_double {  
        costs[j + 1] = newCost as c_float;  
    }  
    ...  
}
```

0 means...

- 
- 0
 - 0.0
 - False
 - Null pointer
 - End of string
 - Sadness
 - Frustration

Benefits?

Safety



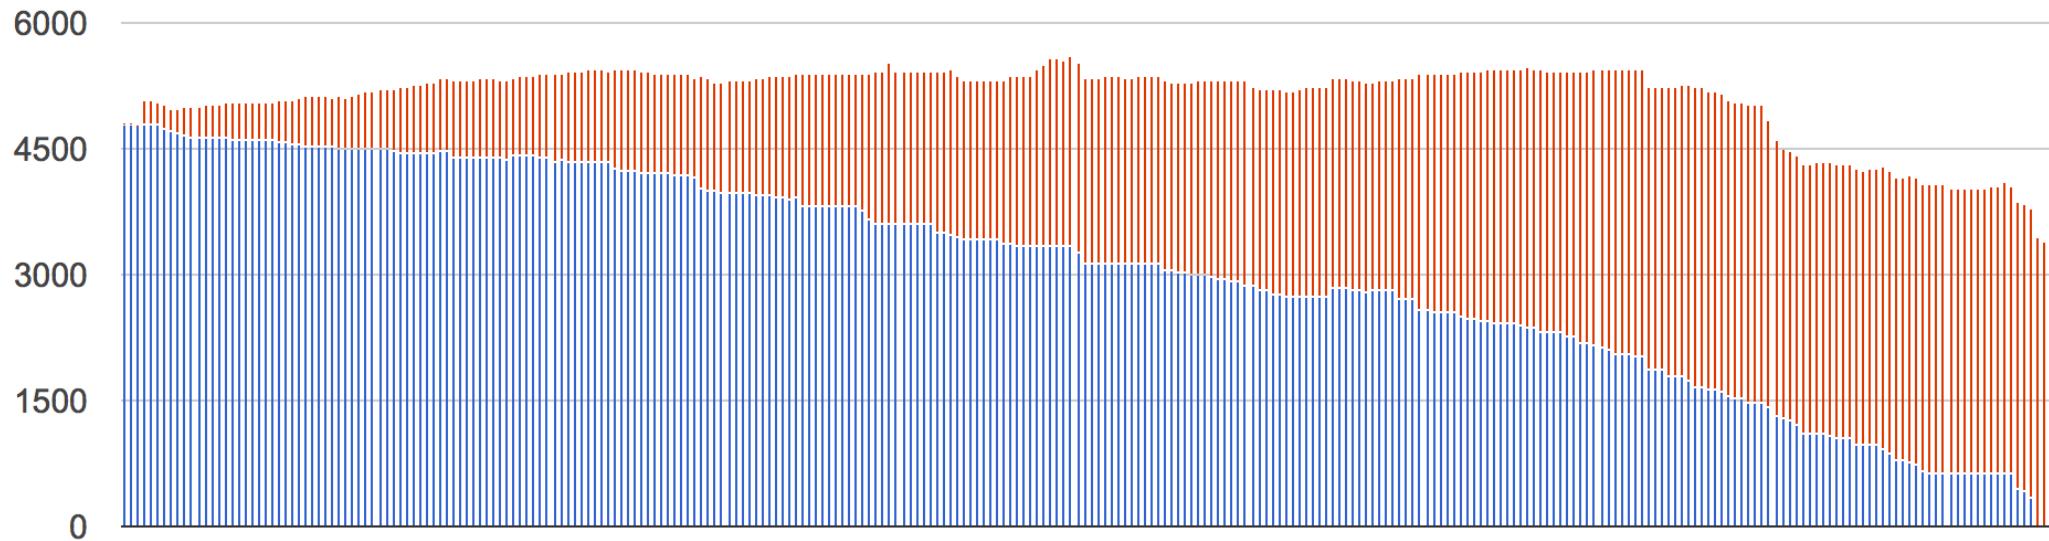
Clarity

Less code

```
#define ZOPFLI_APPEND_DATA(value, data, size) {\  
    if (!((*size) & ((*size) - 1))) {\\  
        /*double alloc size if it's a power of two*/\  
        (*data) = (*size) == 0 ? malloc(sizeof(**data))\  
                            : realloc((*data), (*size) * 2 *  
sizeof(**data));\  
    }\  
    (*data)[(*size)] = (value);\  
    (*size)++;\  
}
```

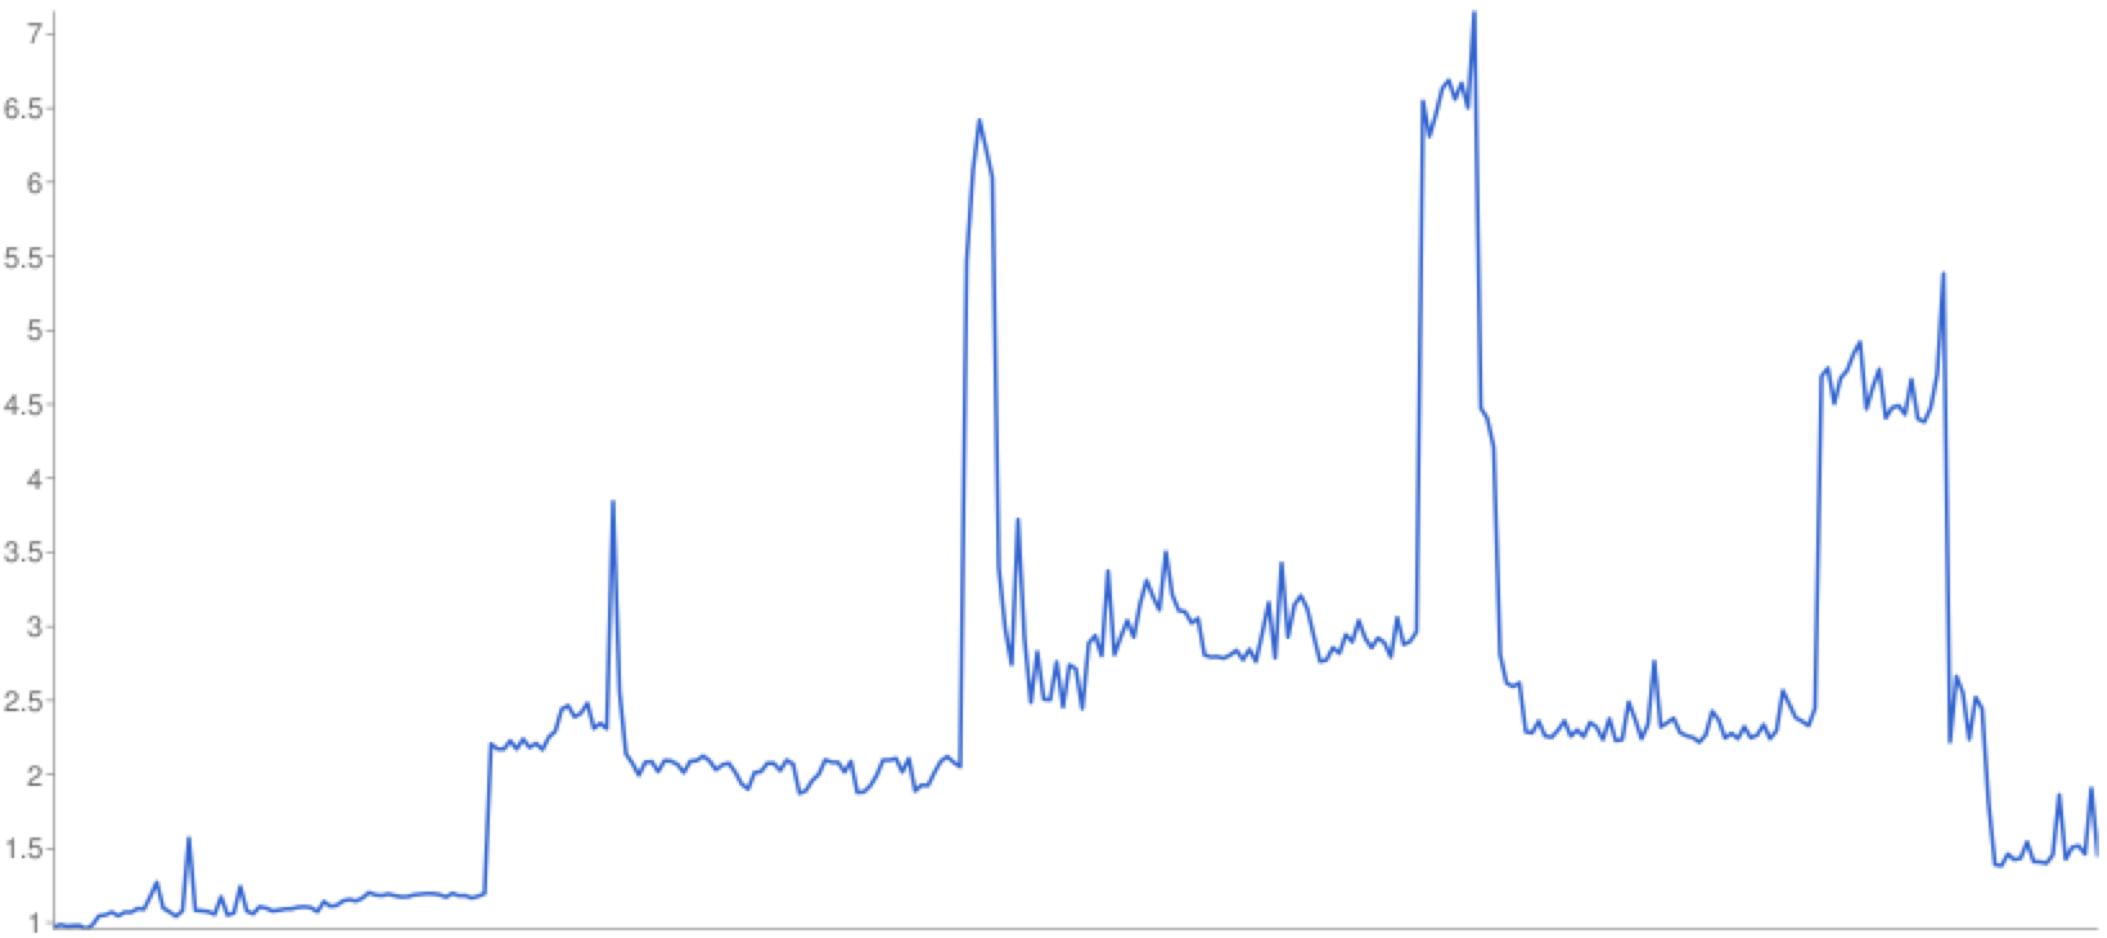
```
#define ZOPFLI_APPEND_DATA(value, data, size) {\  
    if (!((*size) & ((*size) - 1))) {\\  
        /*double alloc size if it's a power of two*/\  
        (*data) = (*size) == 0 ? malloc(sizeof(**data))\  
                               : realloc((*data), (*size) * 2 *  
sizeof(**data));\  
    }\  
    (*data)[(*size)] = (value);\  
    (*size)++;\  
}
```

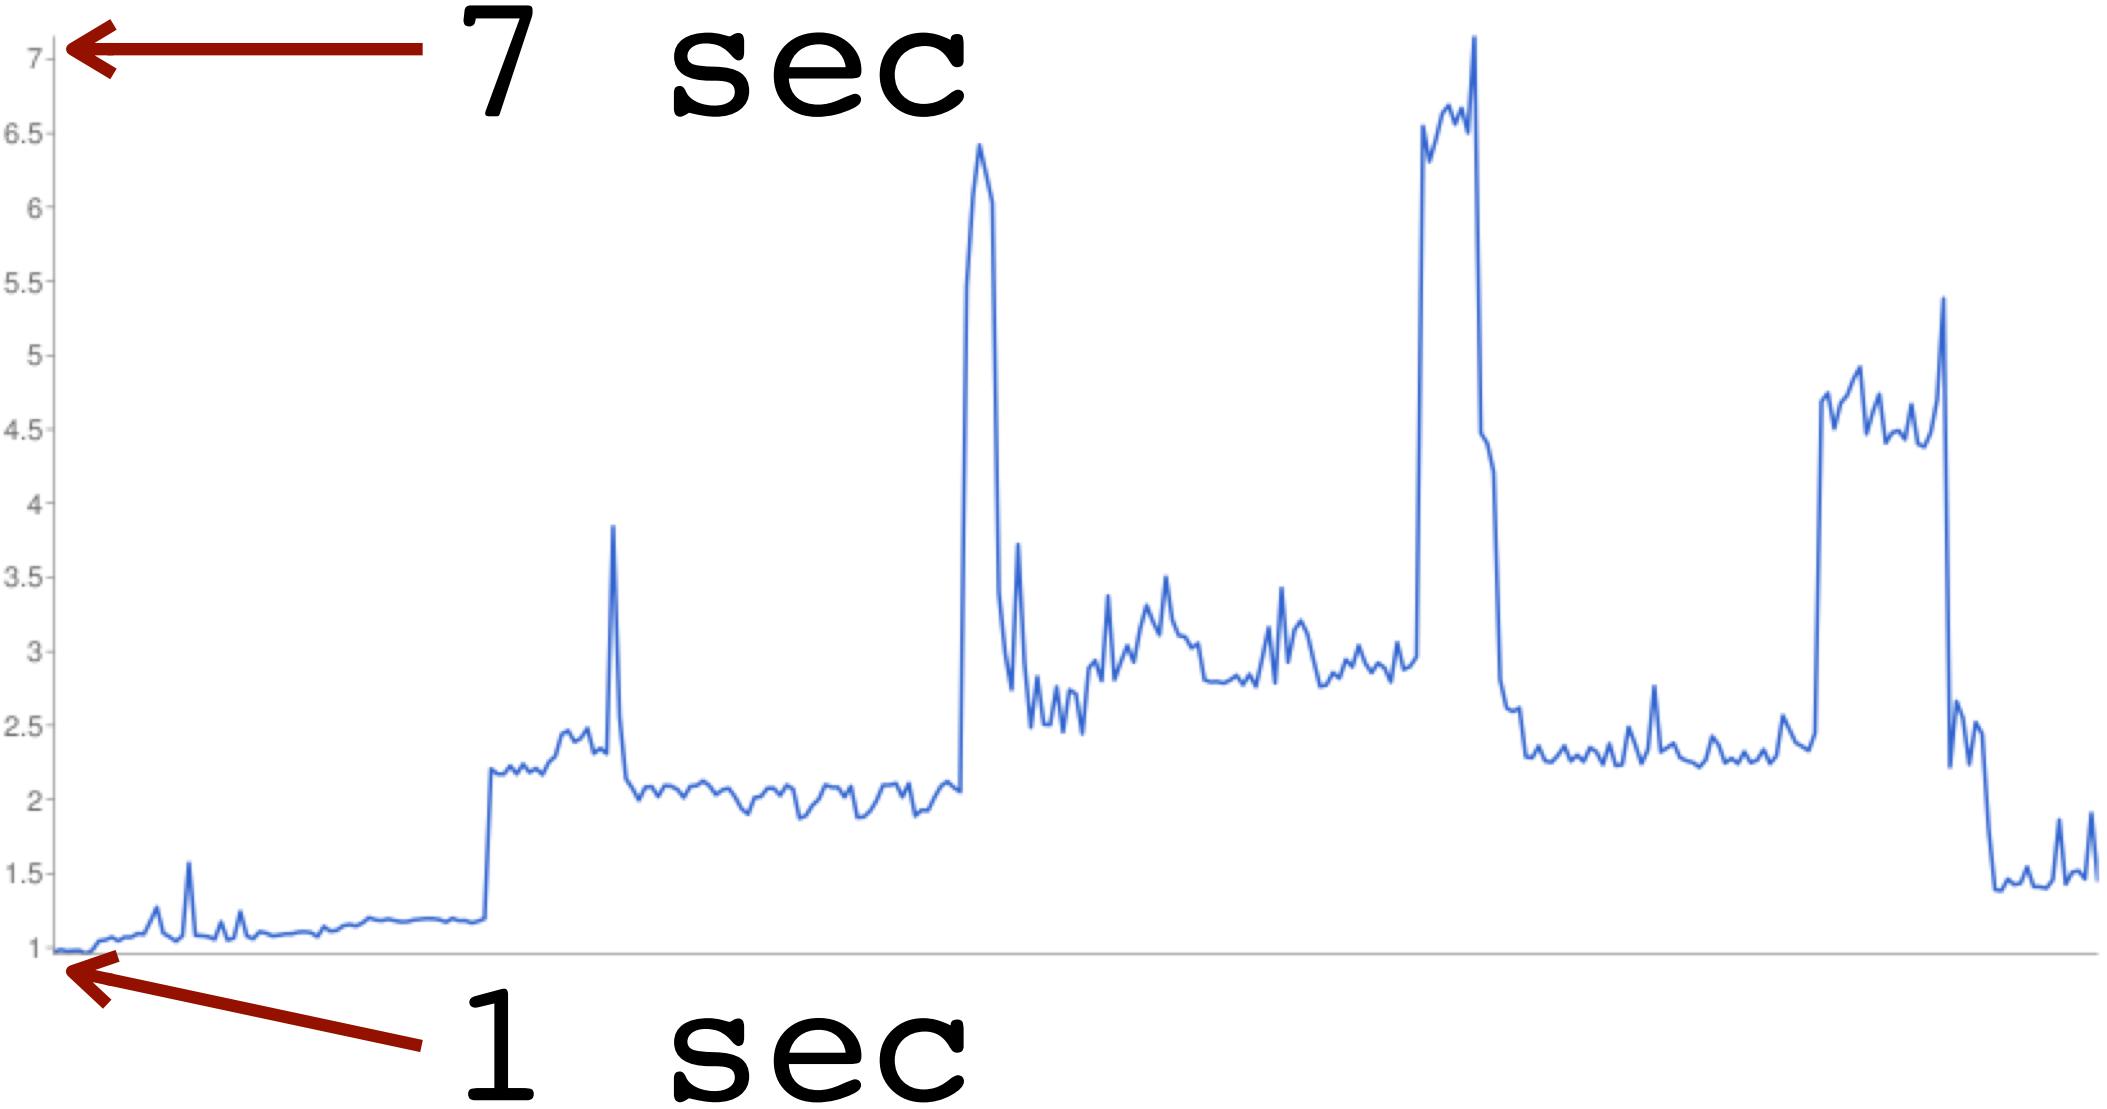
data.push(value)



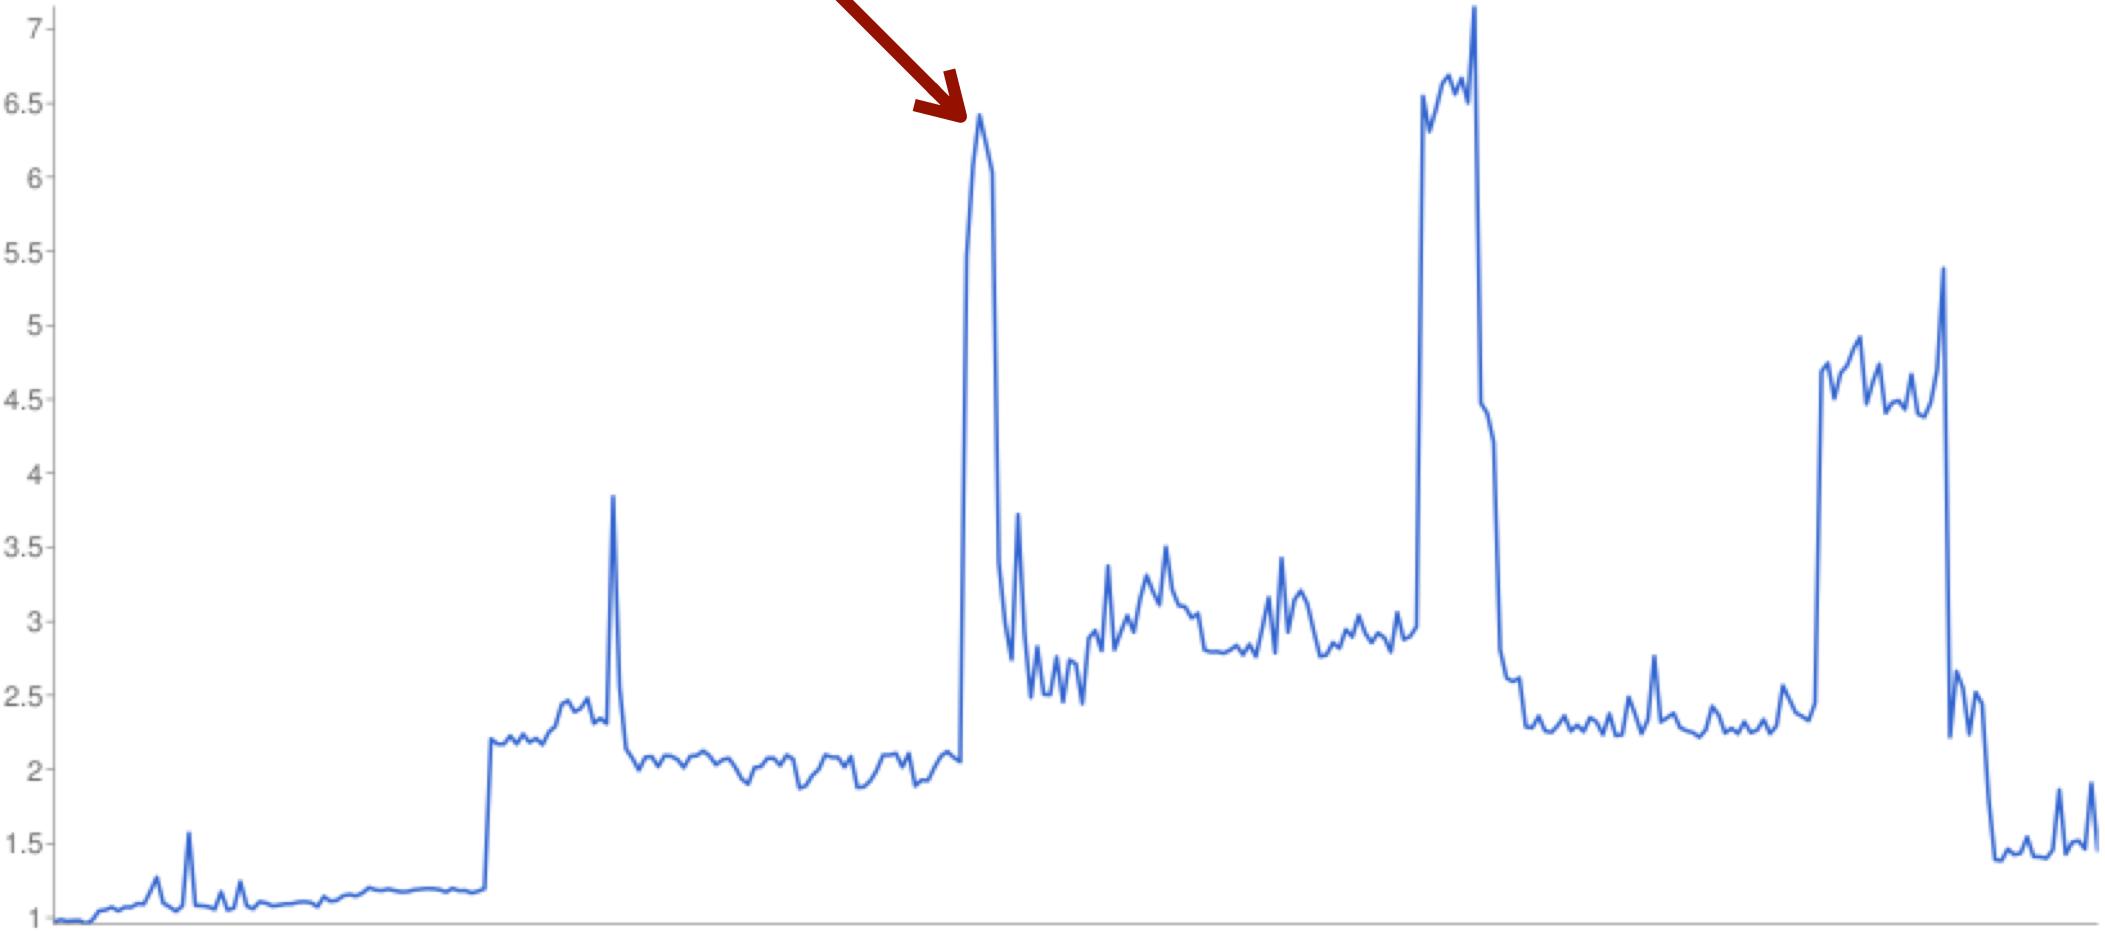
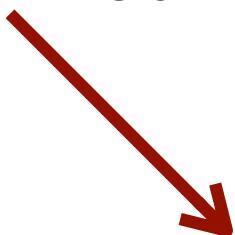
4777 LOC in C ->
3399 LOC in Rust =
71%

Performance

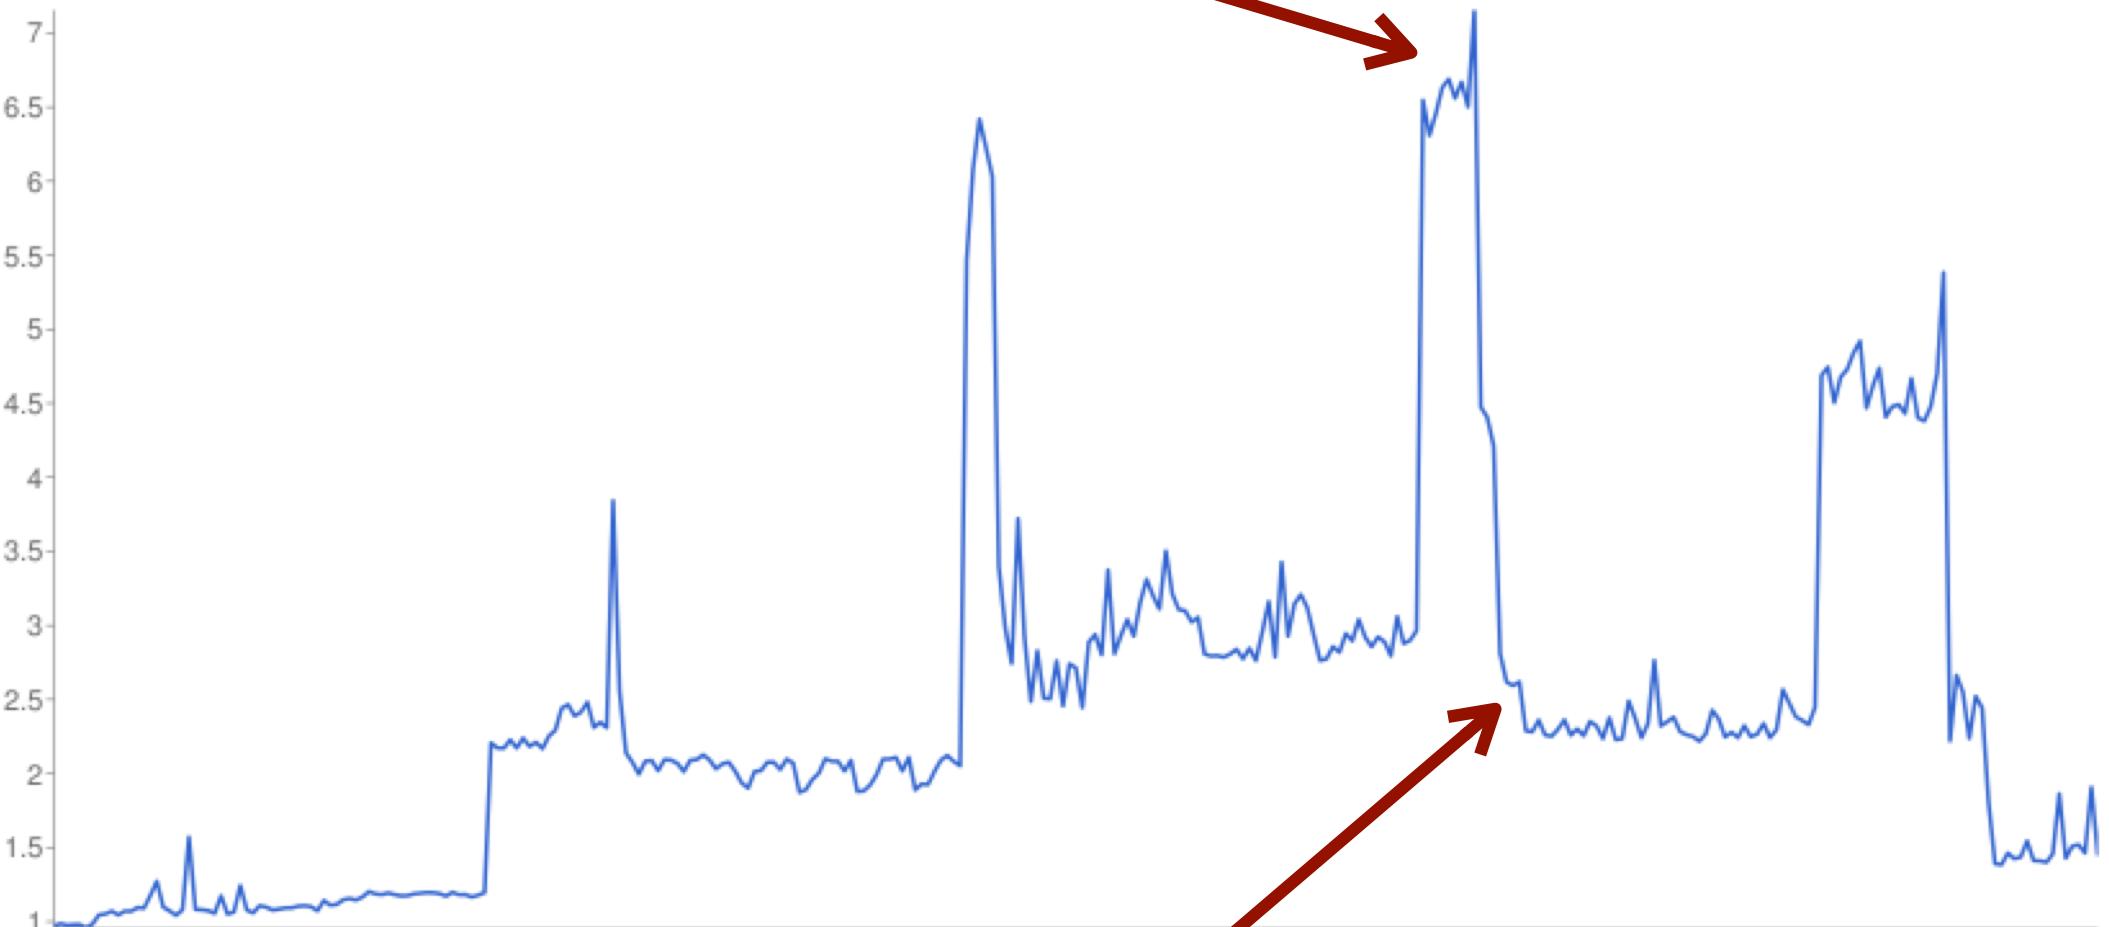




cloning/creating objects
instead of modifying

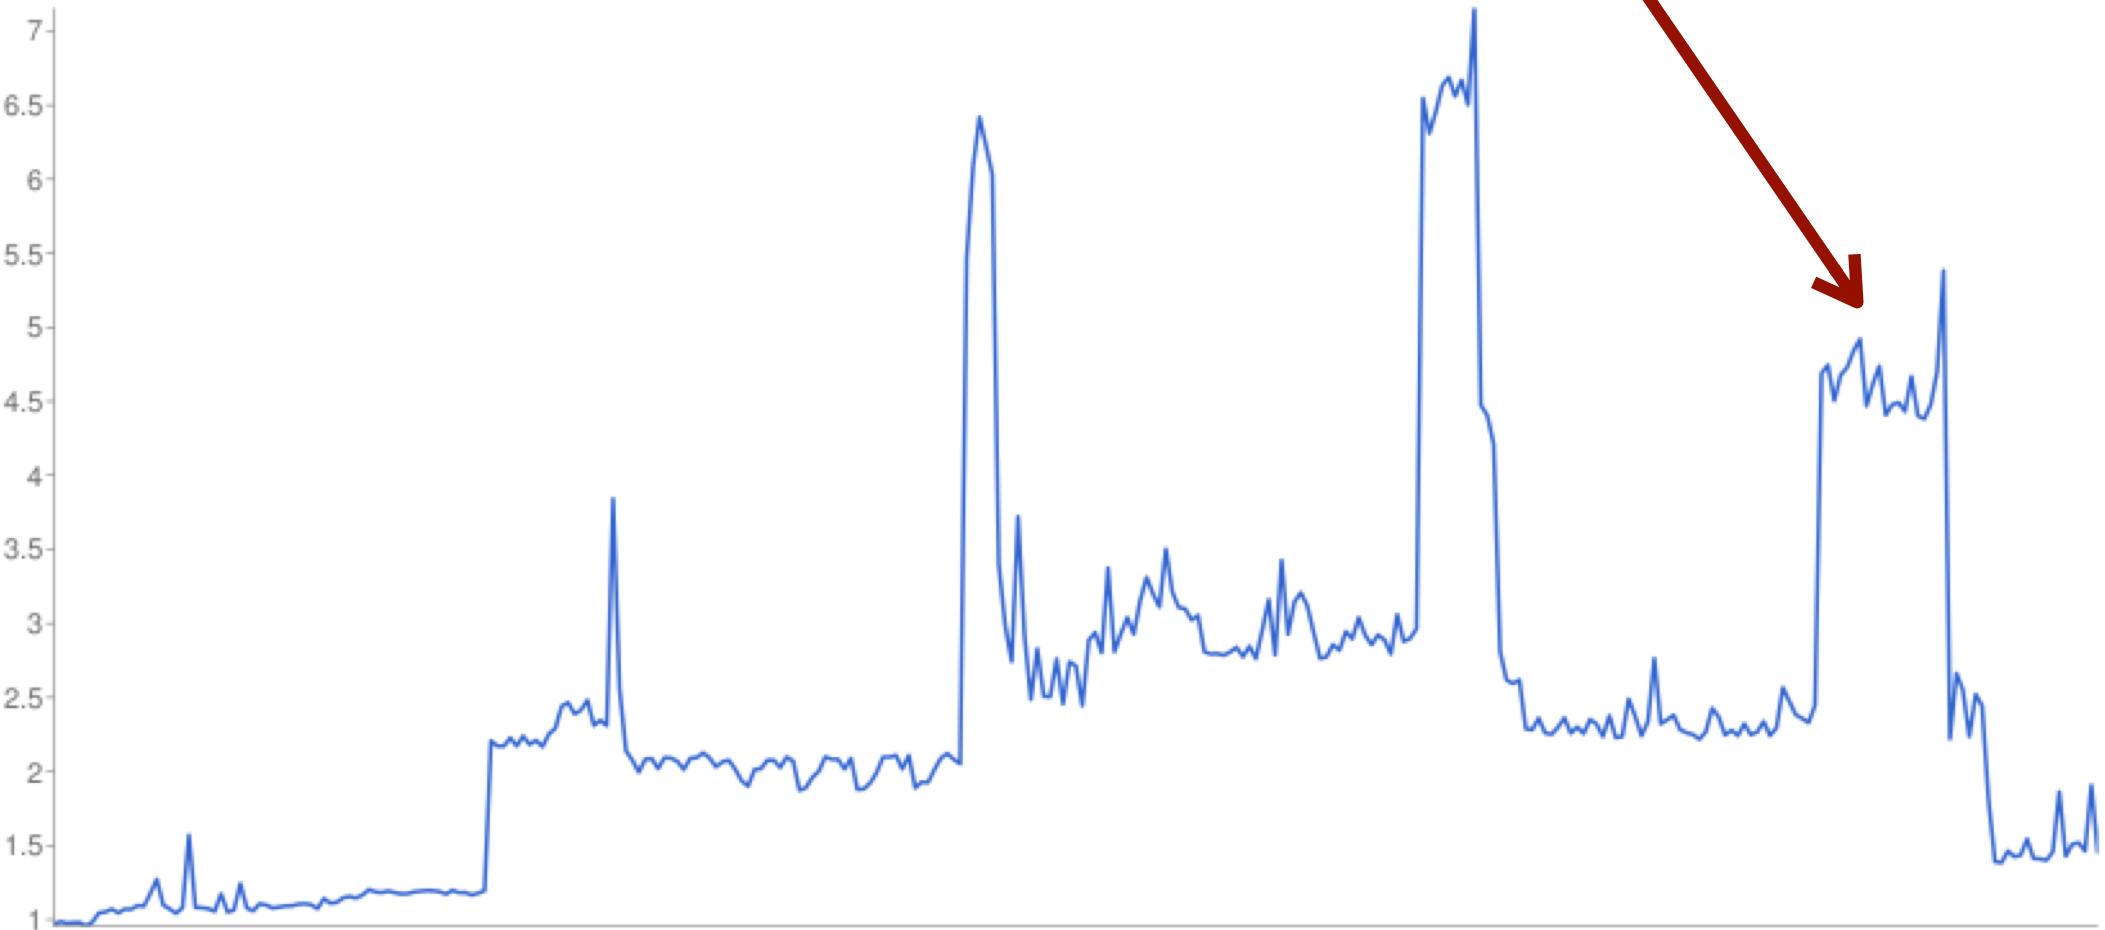


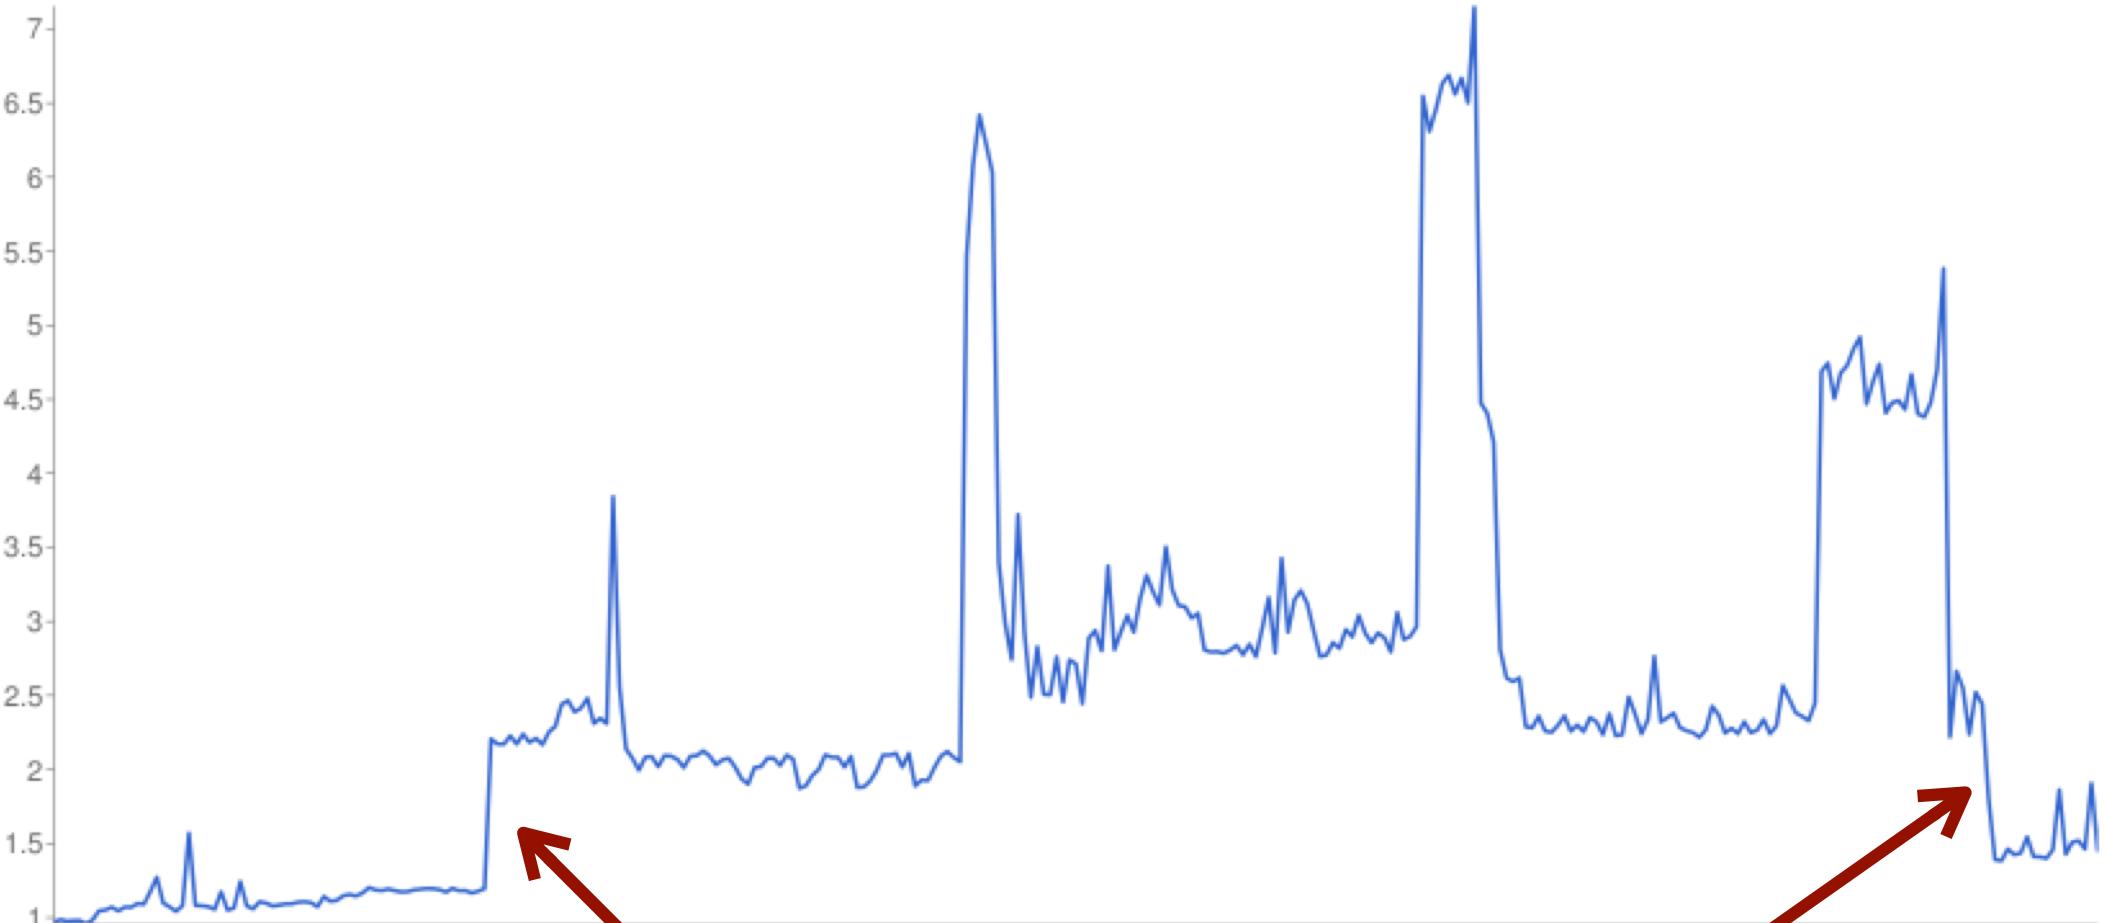
converting lots of
ptrs to Vecs

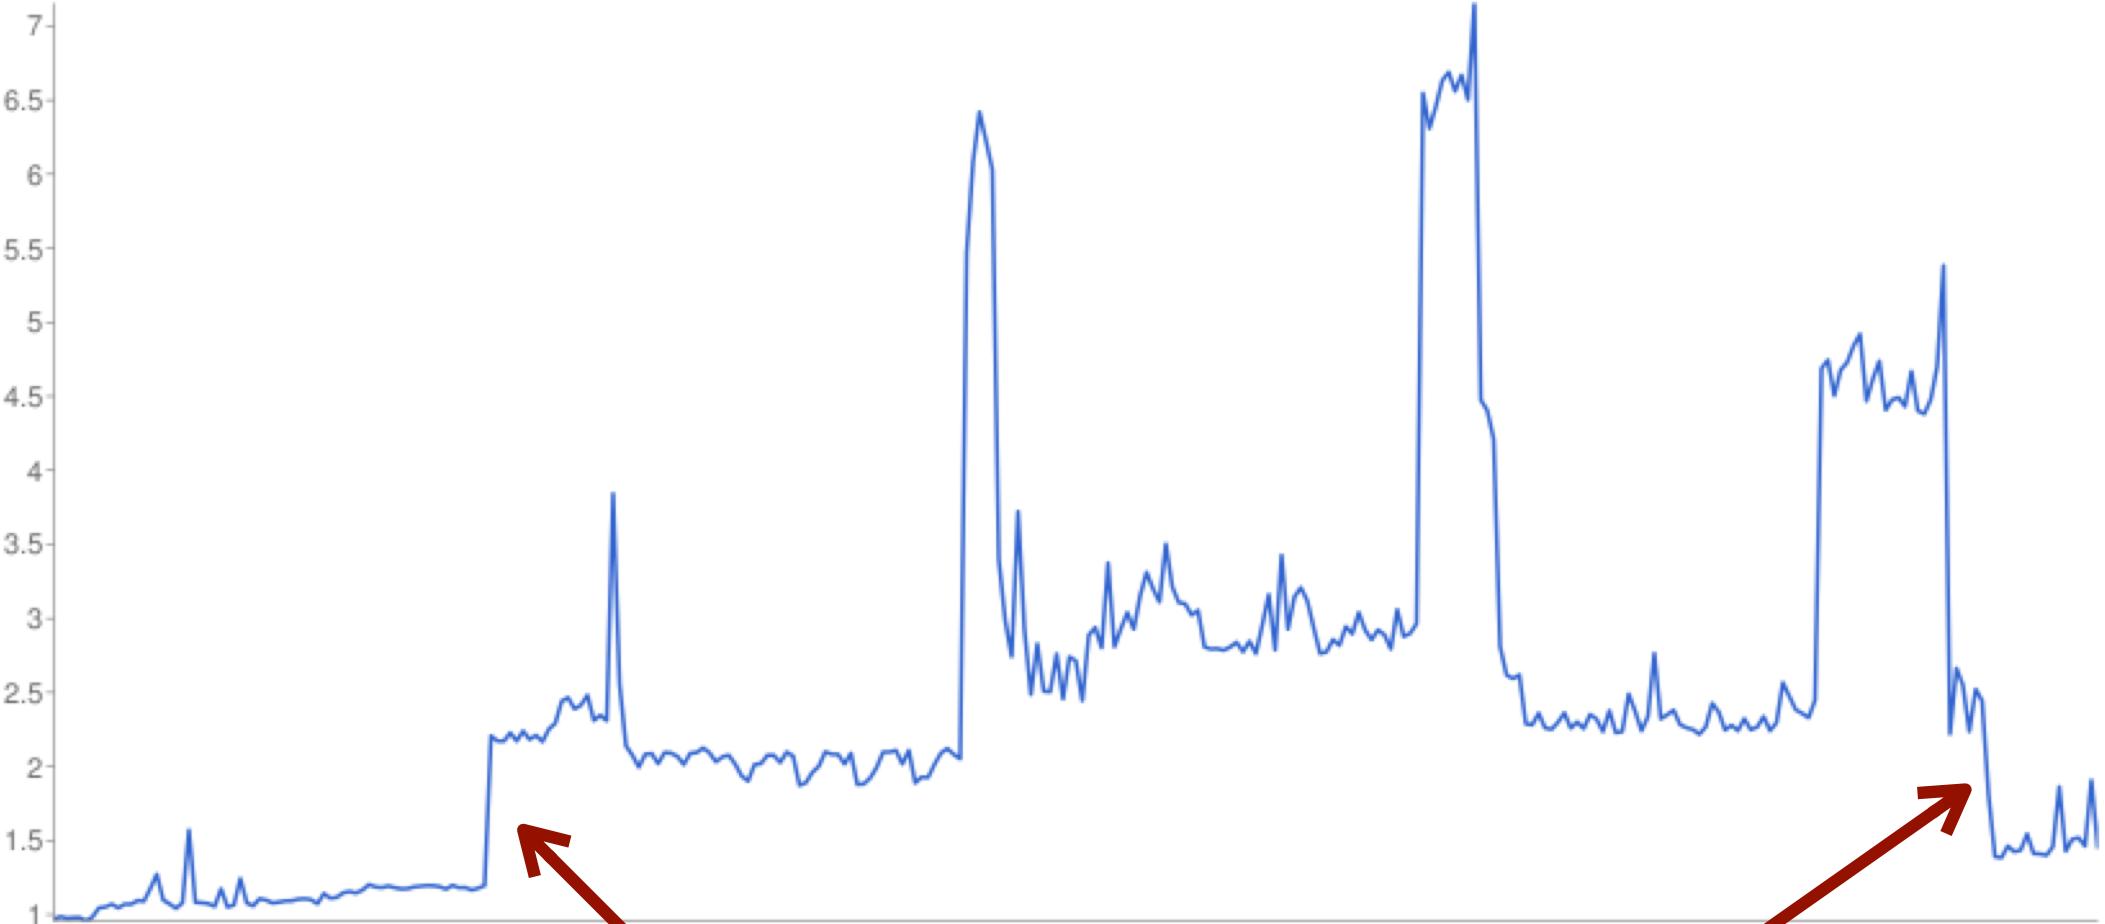


everything converted
to Vecs/slices

copying lots of doubles

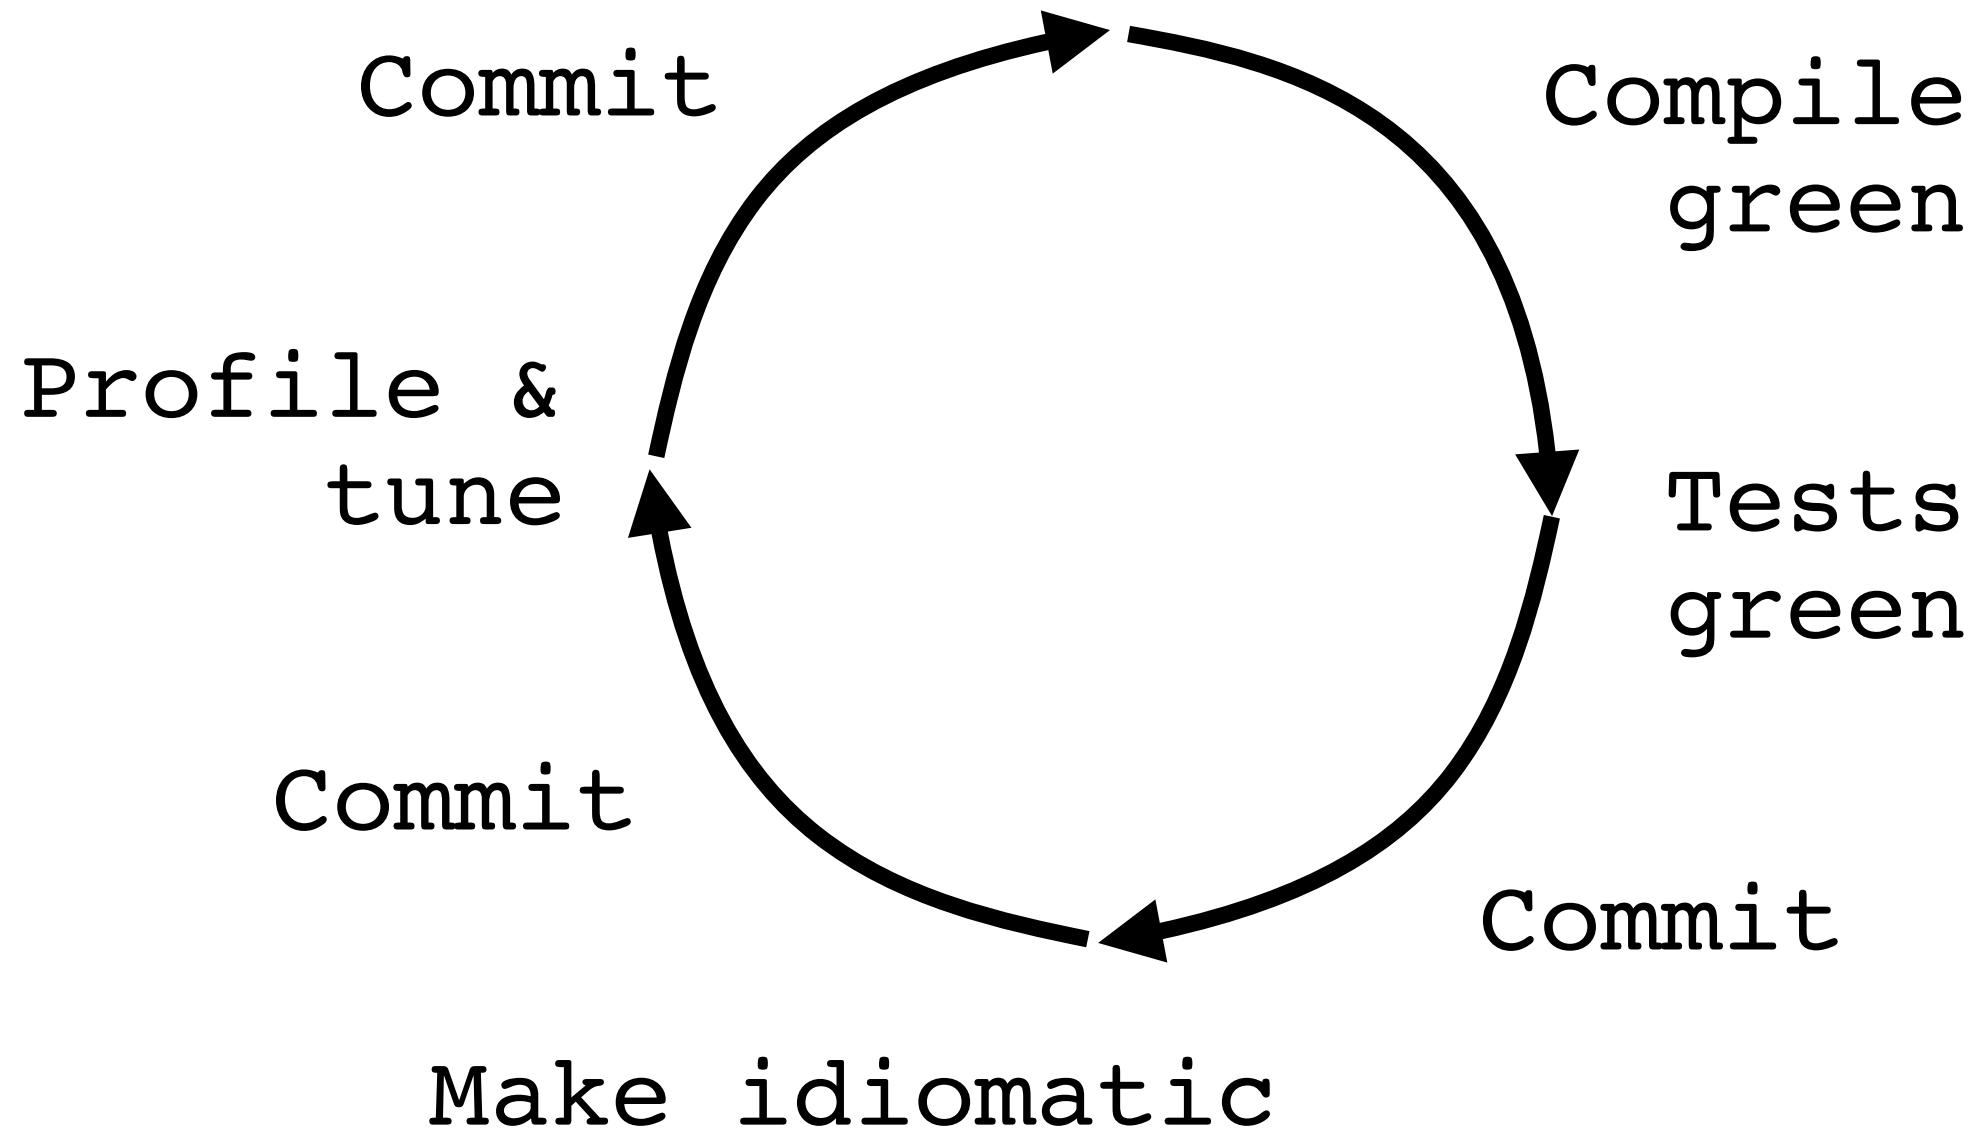






“caching” to
local vars :(

Move C to Rust



C-like Rust is
slower than
idiomatic Rust?

Future work

- Remove all `unsafe`
- Use more iterators
- Stream input/output
- Refactor forever

- \ (ツ) / -

my point:

- Incremental rewrites from C to Rust are possible.
- Have reasons for rewriting and measure progress against the reasons.

References (is.gd/c_rust)

- Repo of my code
- These slides
- FFI chapter in The Rust Programming Language book
- Rust FFI Omnibus
- Working Effectively with Legacy Code by Michael Feathers

Thank You

Carol (Nichols || Goulding)
@carols10cents